

# XiangShan: An Open-Source High-Performance RISC-V Processor and Infrastructure for Architecture Research

---

*The XiangShan Team*

ASPLOS'24@San Diego, USA

April 27, 2024

# What we will cover in this tutorial

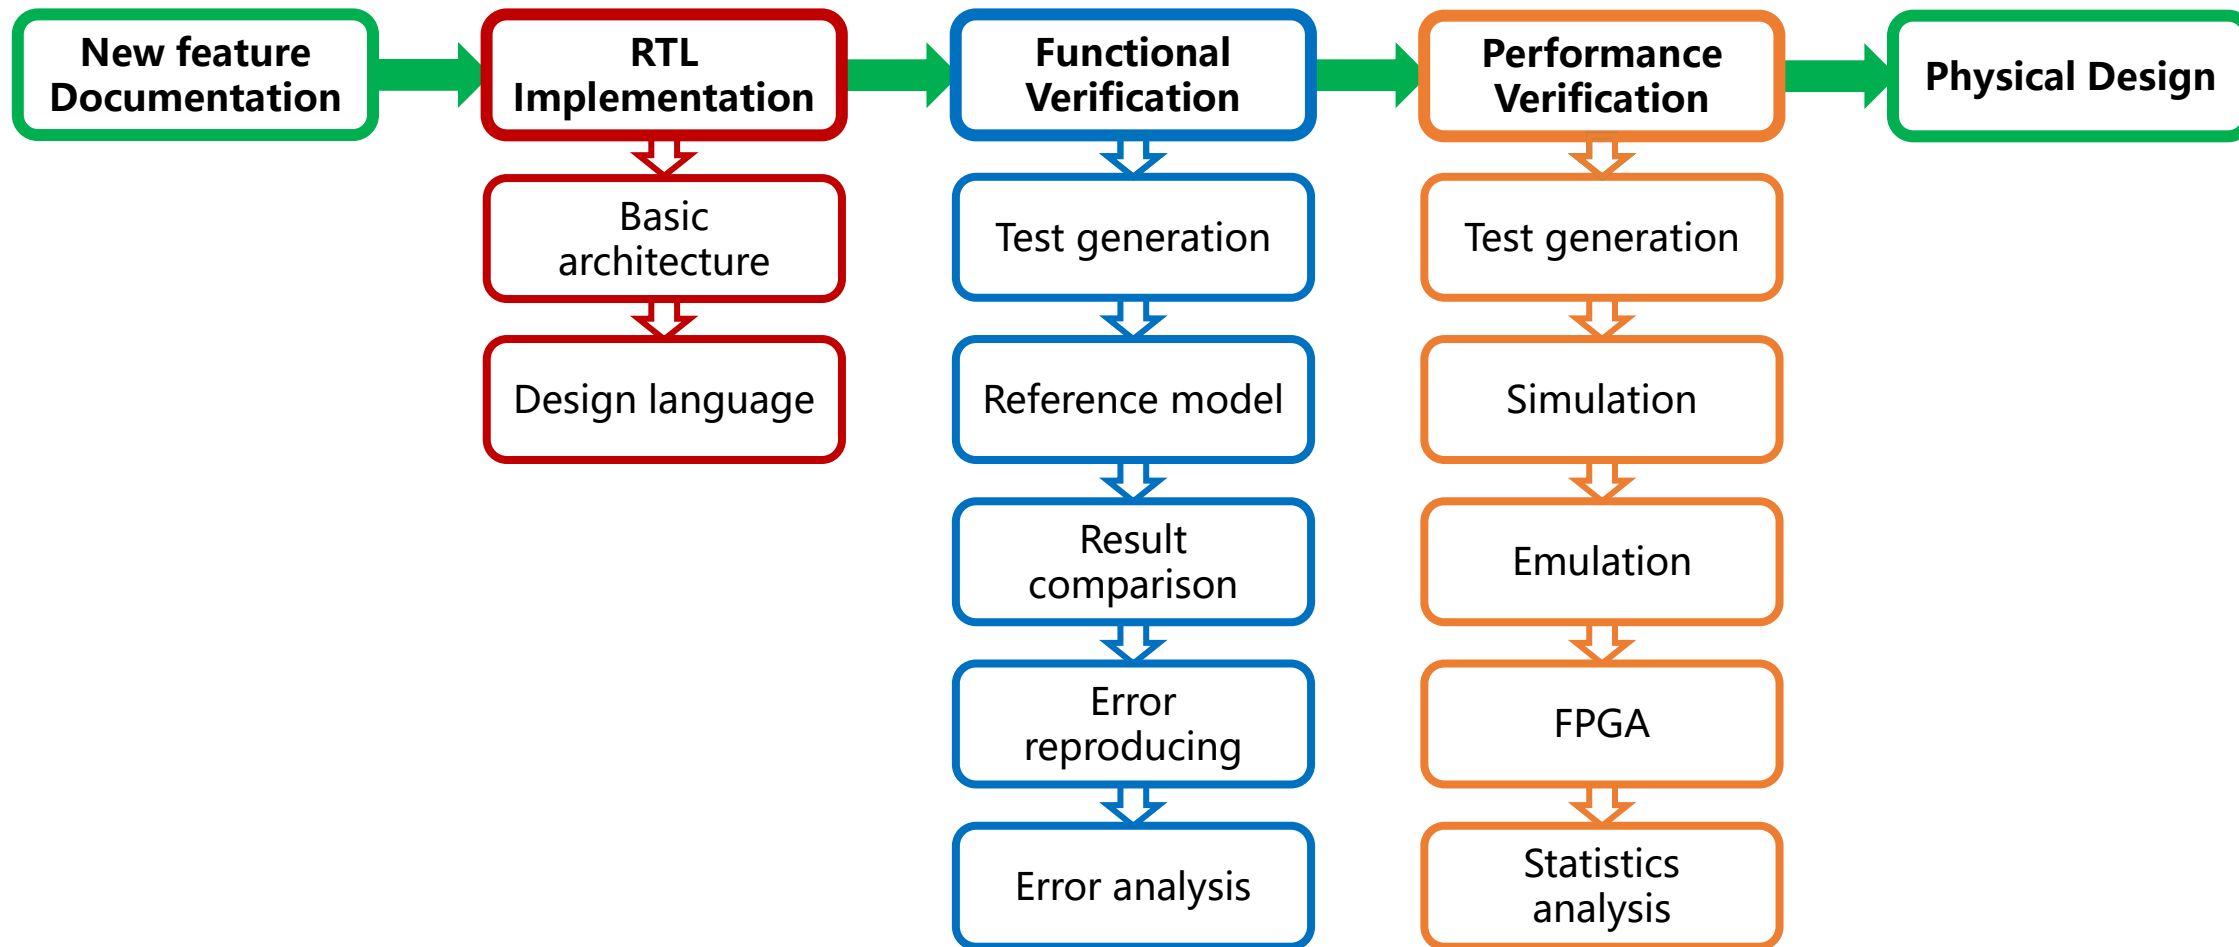
- Introduce of XiangShan (30 minutes)
- CPU Microarchitecture (60 minutes)
  - Design and implementation – How to implement novel ideas on XiangShan
  - Frontend: branch prediction and instruction fetch
  - Backend: out-of-order scheduler, execution units
  - Load/Store Unit: LSQ, pipelines, TLBs, data caches
  - L2/L3 caches and prefetchers
- **Development workflows (60 minutes)**
  - Introduction of their usages – How to develop on XiangShan with MinJie
  - Simulation and Debugging
  - Research Demo

# Tutorial@ASPLOS'24 Schedule

Time (PM)	Topic
1:30 - 2:00	Introduction of the XiangShan Project
2:05 - 3:00	Microarchitecture Design and Implementation
	Afternoon Break
3:30 – 5:00	<b>Hands-on Development &amp; Discussions</b>

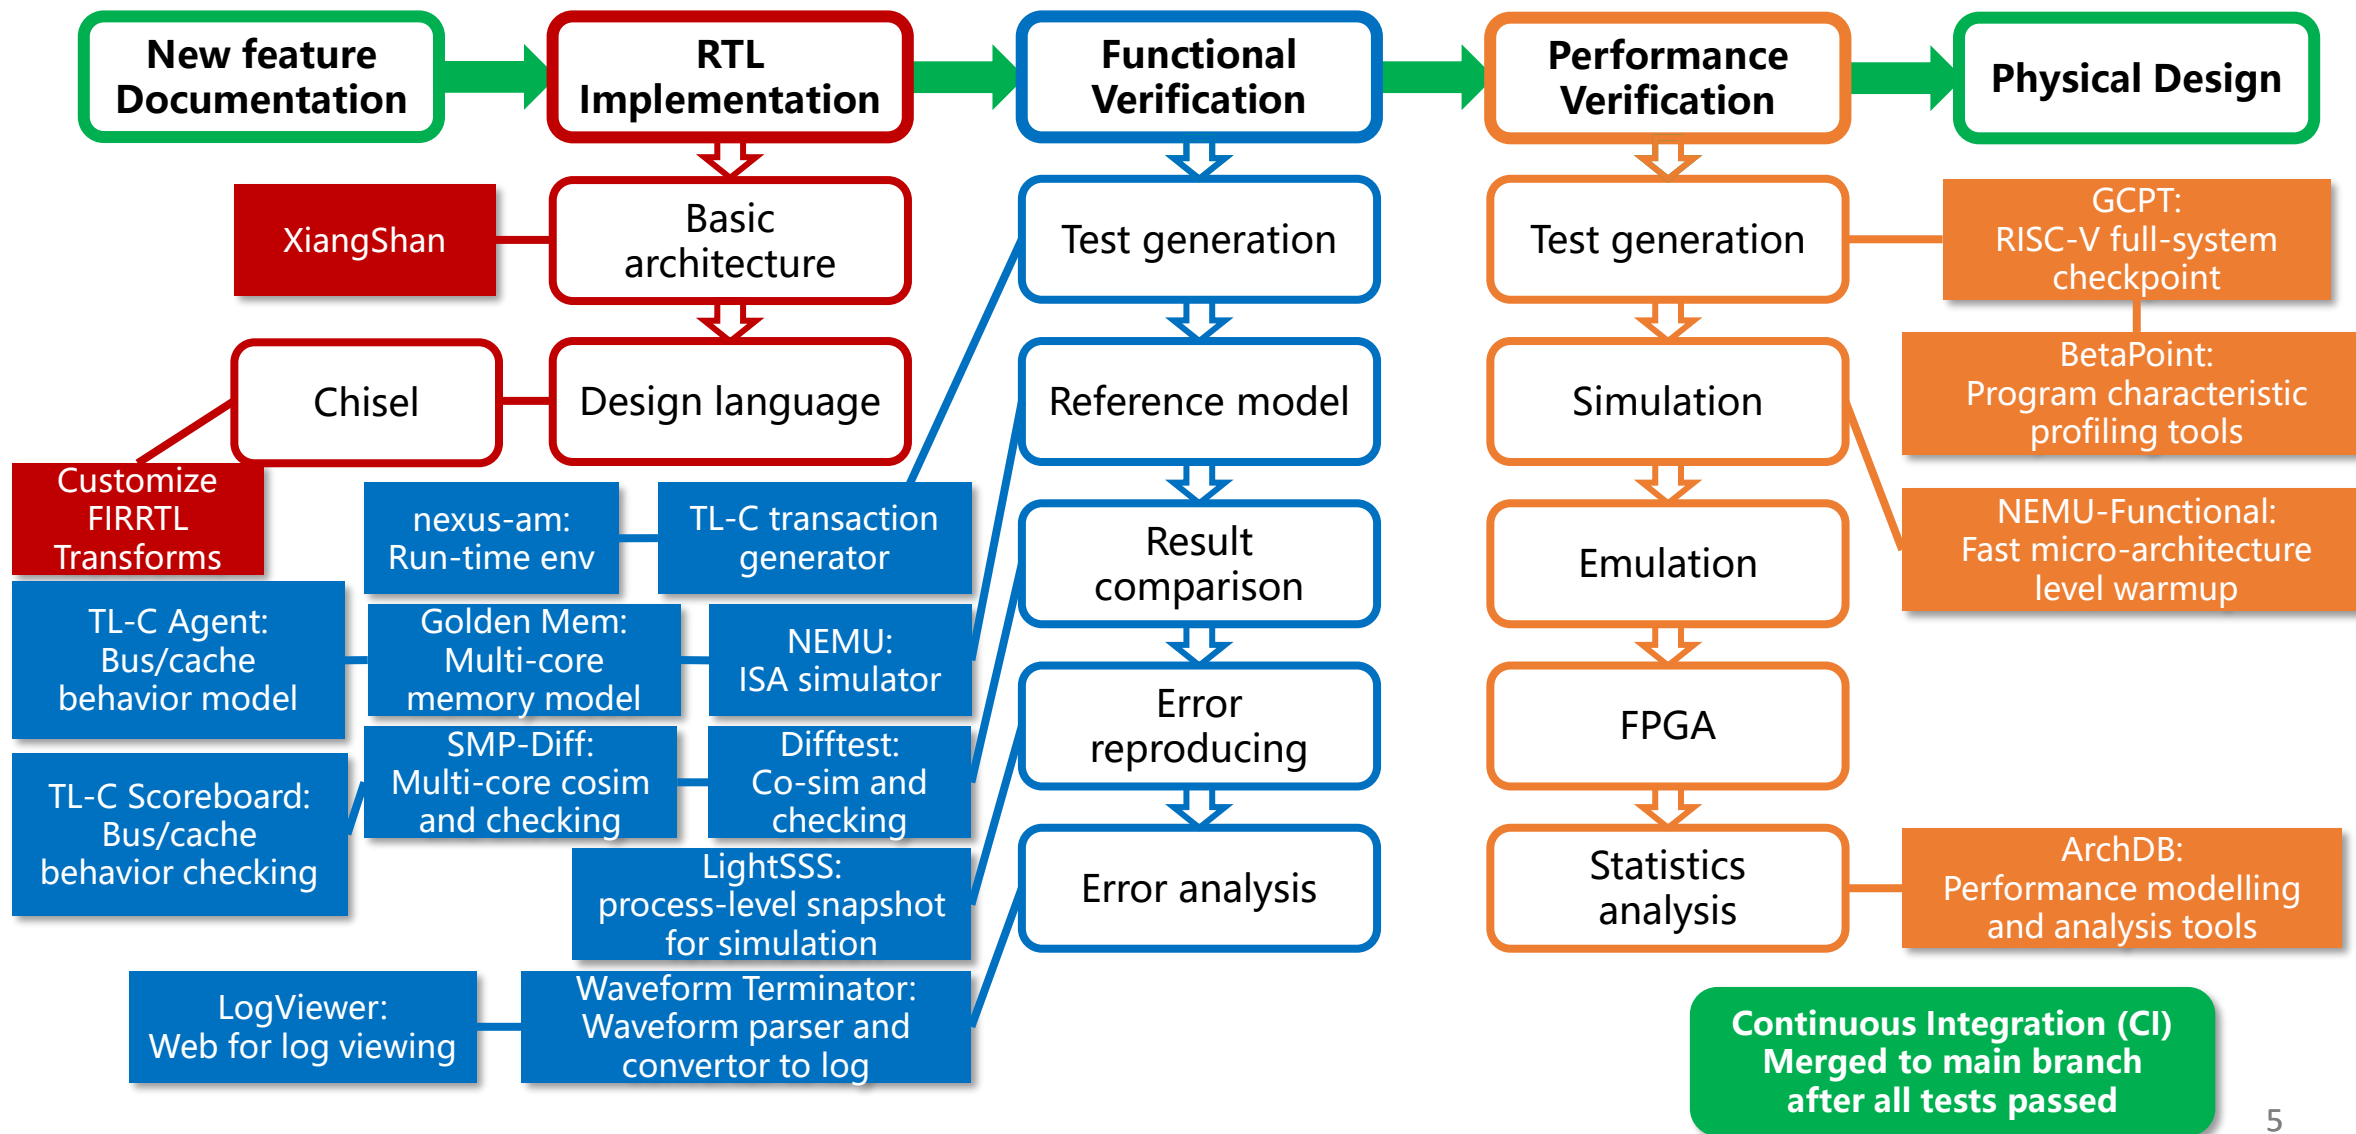


# Roadmap: MinJie Development Flows and Tools





# Roadmap: MinJie Development Flows and Tools





## In this tutorial, we will

- Provide access to cloud servers with you
- Prepare the XiangShan development environment for you
- Go through the development workflows including:

### Simulation

- Environment
- RTL Generation
- RTL Simulation
- .....

### Func. Verification

- Nexus-AM
- NEMU
- Dfftest
- .....

### Perf. Verification

- Perf. counter
- Constin
- Checkpoint
- .....

### Research Demo

- Oracle BP
- Top-Down
- .....

- ***Required: machines with networking and SSH***



# Demo Instructions

- **Interactive shell commands** are highlighted in **brown box** with prefix \$

```
$ echo "Hello, XiangShan"
```

```
$ echo "Have a nice day"
```

- **Description and notes** are presented in **grey box** with prefix #

```
# Please prepare a laptop with an SSH client.
```

```
# Next, let's start the demo session !
```

## Let's Start!

# Prerequisites

- Login to the provided cloud server
  - Windows (Windows Terminal / PowerShell is recommended)

```
PS > ssh guest@47.254.126.201  
# Password: asplos2024
```

- Linux / Mac

```
$ ssh guest@47.254.126.201  
# Password: asplos2024
```

For **offline users**, please refer to <https://github.com/OpenXiangShan/xs-env/tree/tutorial-new>





# Prerequisites

*# Copy tutorial environment to your dir based on your name*

```
$ cp -r /opt/xs-env ~/<YOUR_NAME>
```

*# Enter your dir*

```
$ cd ~/<YOUR_NAME>
```

*# Set up environment variables*

*# **DO IT AGAIN** when opening a new terminal*

```
$ source env.sh
```

```
# SET XS_PROJECT_ROOT:                /home/guest/YOUR_NAME
# SET NOOP_HOME (XiangShan RTL Home):  $XS_PROJECT_ROOT/XiangShan
# SET NEMU_HOME:                       $XS_PROJECT_ROOT/NEMU
# SET AM_HOME:                         $XS_PROJECT_ROOT/nexus-am
```



# Prerequisites

## *# Project Structure*

```
$ tree -d -L 1
```

```
.
├── env-scripts    Useful scripts
├── NEMU           NEMU
├── XiangShan      XiangShan processor
├── nexus-am       Abstract machine
└── tutorial       Scripts and patches for tutorial
```

## *# Enter XiangShan directory*

```
$ cd XiangShan
```



# Chisel Compilation



- Compile RTL and build simulator with Verilator

```
$ make emu CONFIG=MinimalConfig MFC=1 -j8
```

*Options:*

<i>CONFIG=MinimalConfig</i>	<i>Configuration of XiangShan</i>
<i>MFC=1</i>	<i>Enable MLIR Firrtl Compiler</i>
<i>// EMU_THREADS=2</i>	<i>Simulation threads</i>
<i>// EMU_TRACE=1</i>	<i>Enable waveform dump</i>
<i>// WITH_DRAMSIM=1</i>	<i>Enable DRAMSim3 for DRAM simulation</i>
<i>// WITH_CHISELDB = 1</i>	<i>Enable ChiselDB feature</i>
<i>// WITH_CONSTANTIN = 1</i>	<i>Enable Constantin feature</i>

- Compilation might take **10 mins~**



# Open Another Terminal

*# login to the cloud server again*

```
$ ssh guest@47.254.126.201
```

```
# Password: asplos2024
```

*# Come back to XiangShan directory*

```
$ cd ~/<YOUR_NAME> && source env.sh && cd XiangShan
```

# Run RTL Simulation with Verilator

- After building, we can run the simulator

*# run pre-built simulator*

```
$ ../tutorial/emu -i ../tutorial/hello.bin --no-diff 2>hello.err
```

*# Some key options:*

*-i*

*Workload to run*

*-C / -I / -W*

*Max cycles / Max Insts / Warmup Insts*

*--diff=PATH / --no-diff*

*Compare with NEMU for difftest / disable difftest*

**Great!**

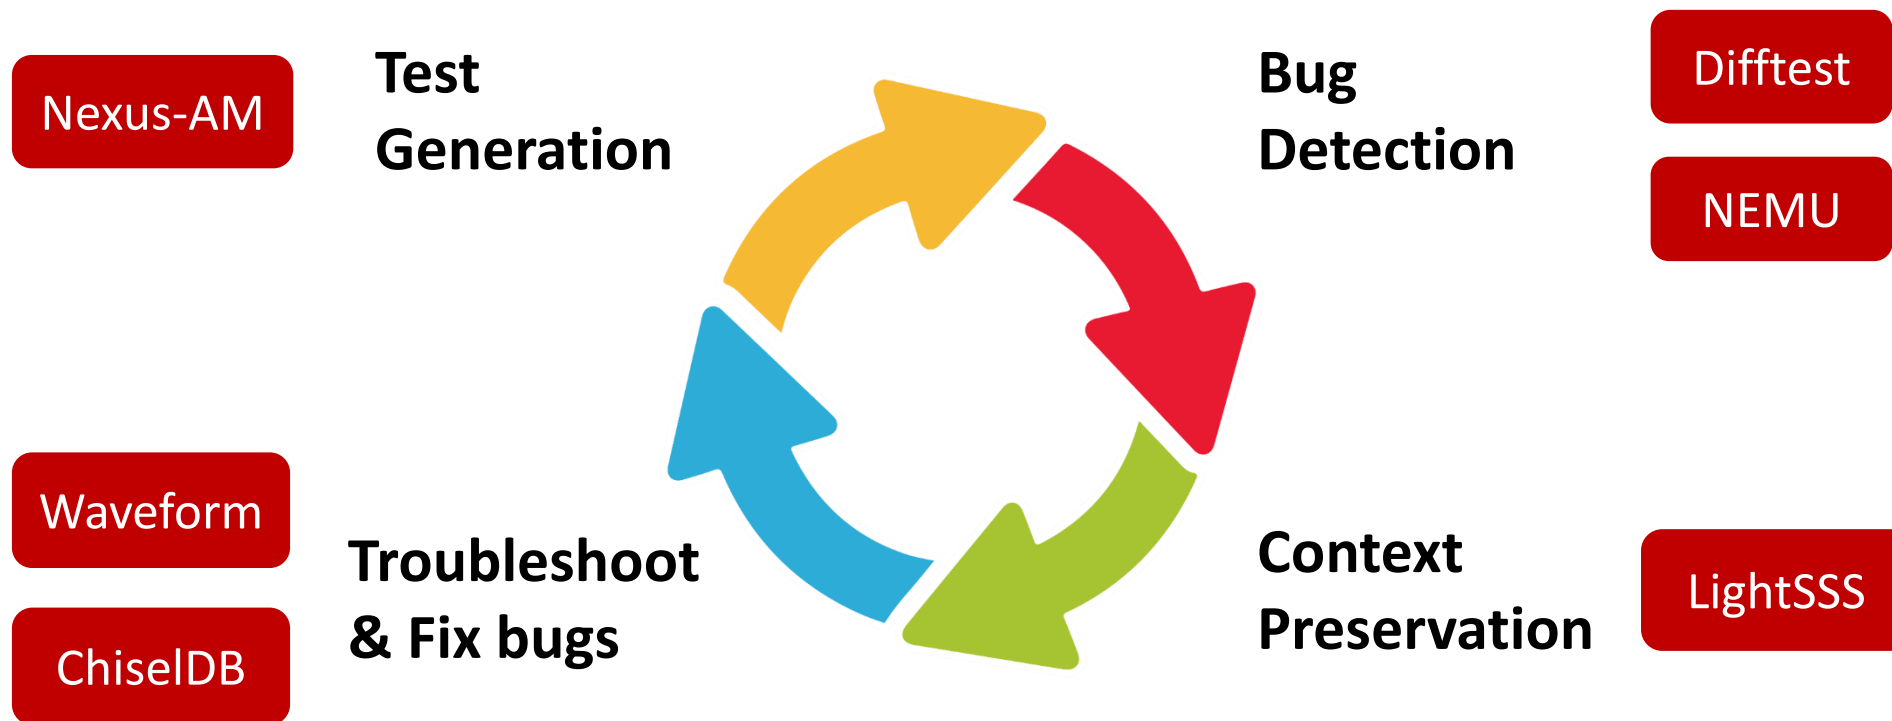
**We have learned the basic simulation process of Xiangshan.**

Once we get the RTL simulator ready, what's next?

How can we

- **generate desired workloads**
- **find and fix functional bugs**
- **do performance analysis**
- **do research on micro-architecture**

# MinJie Functional Verification



# Bare Metal Runtime: Nexus-AM

- **Purpose**
  - Generate workloads **agilely** without an OS
  - Provide runtime framework for **bare metal machines** like XiangShan
- **We present a bare metal runtime called Nexus-AM**
  - Light-weight, easy to use
  - Implement basic **system call** interfaces and exception handlers
  - Support multiple **ISAs and configurations**





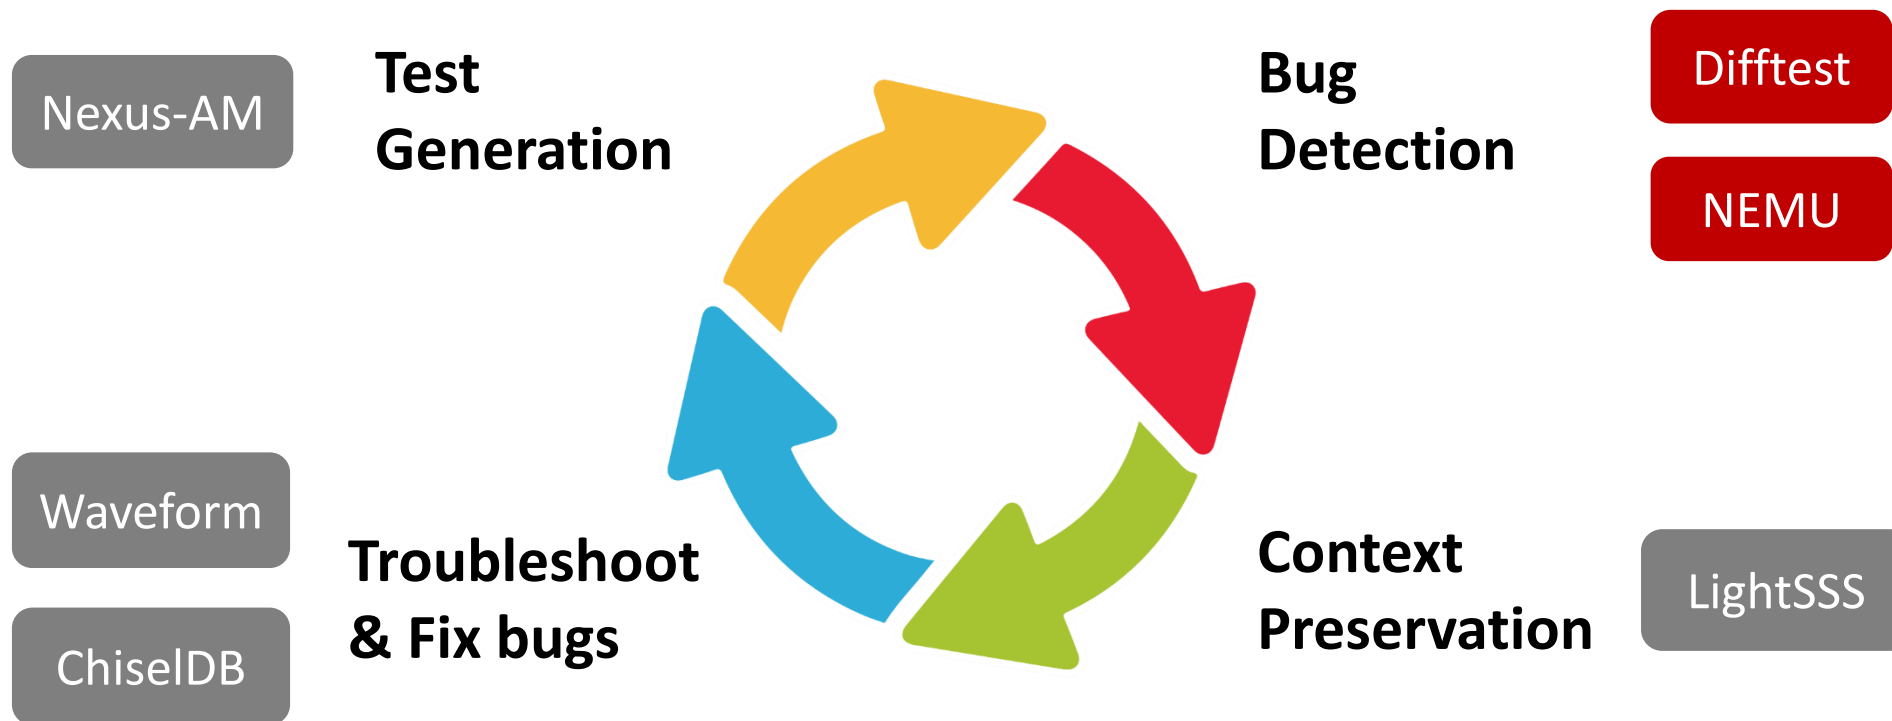
# Nexus-AM

- Hands-on: build Coremark workload
- Showcase

```
$ cd ~/<YOUR_NAME> && source env.sh && cd tutorial  
$ bash build-am.sh
```

```
# cd $AM_HOME/apps/coremark  
# make ARCH=riscv64-xs  
# ls -l build  
#   coremark-riscv64-xs.bin    Binary image of the program  
#   coremark-riscv64-xs.elf    ELF of the program  
#   coremark-riscv64-xs.txt    Disassembly of the program
```

# MinJie Functional Verification





# ISA REF: NEMU

- **Purpose**

- Golden model for verification
- Simple yet high performance

- **We present NEMU**

- An **instruction set simulator**, similar to Spike
- Through **optimization techniques**, achieve performance similar to QEMU.
- **A set of APIs** is provided to assist XiangShan in comparing and verifying the **architectural states**

- Showcase

```
$ bash build-nemu.sh
```

```
# cd $NEMU_HOME
```

```
# make clean
```

```
# make riscv64-xs_defconfig
```

```
# make -j build NEMU as the bare metal machine, can run the Coremark from the previous step
```

```
# make clean-softwarefloat
```

```
# make riscv64-xs-ref_defconfig
```

```
# make build NEMU as the reference model for XiangShan
```

- **Hands-on:** run Coremark workload on NEMU
- **Showcase**

```
$ bash run-nemu.sh
```

```
# cd $NEMU_HOME
```

```
# ./build/riscv64-nemu-interpretter -b \
```

*run in batch mode, faster*

```
$AM_HOME/apps/coremark/build/coremark-riscv64-xs.bin
```

*set workspace to Coremark*

# ISA co-simulation using Difttest

- **Basic flows**

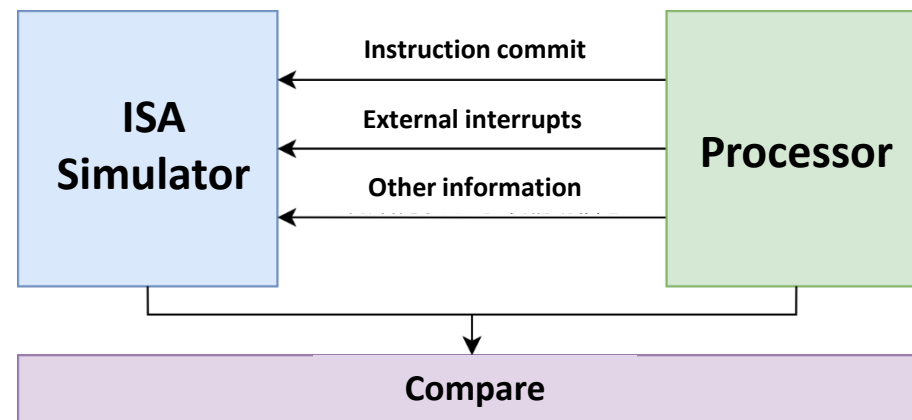
- Instructions commit/other states update
- The simulator executes the same instructions
- Compare the architectural states between DUT and REF
- Abort or continue

- **Provided as verification infrastructure for processors**

- APIs for HDLs such as Chisel/Verilog
- RTL simulators such as Verilator, VCS
- RISC-V ISS such as **NEMU**, **Spike**

- **SMP-Difttest: co-simulation on an SMP processor**

- SMP Linux kernel and multi-threading programs
- Online checking of cache coherency and memory consistency



## Basic architecture of DiftTest

```
while (1) {  
    icnt = cpu_step();  
    nemu_step(icnt);  
    r1s = cpu_getregs();  
    r2s = nemu_getregs();  
    if (r1s != r2s) { abort(); }  
}
```

## Online checking

# DiffTest

- **Hands-on:** run Coremark workload on XiangShan, diffTest with NEMU
- **Showcase**

*# Running Coremark takes about 2 minutes*

```
$ bash run-emu.sh
```

```
# ./emu \      if the simulator is successfully compiled, you can use $NOOP_HOME/build/emu  
-i $AM_HOME/apps/coremark/build/coremark-riscv64-xs.bin \ use coremark.bin  
--diff $NEMU_HOME/build/riscv64-nemu-interpreter-so \      diffTest with NEMU  
2>perf.out      redirect stderr to file
```



# DiffTest

```
./emu -i $AM_HOME/apps/coremark/build/coremark-riscv64-xs.bin --diff $NEMU_HOME/build/riscv64-nemu-interpreter-so
Emu compiled at Mar 24 2023, 20:33:03
Using simulated 32768B flash
[warning]no valid flash bin path, use preset flash instead
The image is /home/guest/xs-env-asplos2023/nexus-am/apps/coremark/build/coremark-riscv64-xs.bin
Using simulated 8192MB RAM
NemuProxy using /home/guest/xs-env-asplos2023/NEMU/build/riscv64-nemu-interpreter-so
The first instruction of core 0 has committed. DiffTest enabled.
[NEMU] Can not find flash image: (null)
[NEMU] Use built-in image instead
[src/device/io/mmio.c:38,add_mmio_map] Add mmio map 'flash' at [0x0000000010000000, 0x00000000100ffffff]
Running CoreMark for 1 iterations
2K performance run parameters for coremark.
CoreMark Size      : 666
Total time (ms)    : 4288
Iterations         : 1
Compiler version   : GCC9.4.0
seedcrc            : 0xe9f5
[0]crclist         : 0xe714
[0]crcmatrix       : 0x1fd7
[0]crcstate        : 0x8e3a
[0]crcfinal        : 0xe714
Finised in 4288 ms.
=====
CoreMark Iterations/Sec 233
Core 0: HIT GOOD TRAP at pc = 0x80001c52
total guest instructions = 398,372
instrCnt = 398,372, cycleCnt = 517,099, IPC = 0.770398
Seed=0 Guest cycle spent: 517,103 (this will be different from cycleCnt if emu loads a snapshot)
Host time spent: 77,559ms
```



# Difftest

- **Hands-on:**

- We injected a bug in a pre-built XiangShan processor
- Now, let's trigger it on the pre-built buggy XiangShan by Difftest

- **Showcase**

```
$ bash run-emu-diff.sh
```

```
# ./emu-bug \  
-i $AM_HOME/apps/coremark/build/coremark-riscv64-xs.bin \  
--diff $NEMU_HOME/build/riscv64-nemu-interpretter-so # difftest with NEMU
```



- Difftest will report the error once failed
  - Dump info like PC, GPRs, etc.
  - Point out the differences

```
===== Commit Group Trace (Core 0) =====
commit group [00]: pc 008000118 cmtcnt 7
commit group [01]: pc 00800012c cmtcnt 6
commit group [02]: pc 00800013b2 cmtcnt 6
commit group [03]: pc 008000124 cmtcnt 7
commit group [04]: pc 008000136 cmtcnt 6
commit group [05]: pc 00800013c0 cmtcnt 6
commit group [06]: pc 00800013da cmtcnt 6
commit group [07]: pc 0080001400 cmtcnt 7
commit group [08]: pc 0080001414 cmtcnt 7
commit group [09]: pc 0080001428 cmtcnt 6
commit group [10]: pc 0080001436 cmtcnt 6
commit group [11]: pc 00800017fe cmtcnt 6
commit group [12]: pc 008000180c cmtcnt 6 <--
commit group [13]: pc 008000139a cmtcnt 6
commit group [14]: pc 0080000124 cmtcnt 7
commit group [15]: pc 0080000154 cmtcnt 6
```

```
===== REF Regs =====
$0: 0x0000000000000000 ra: 0x00000000800013b8 sp: 0x000000008000cf00 gp: 0x0000000000000000
tp: 0x0000000000000000 t0: 0x0000000000000000 t1: 0x00000000800039d0 t2: 0x0000000000000000
s0: 0x000000008000cf10 s1: 0x0000000000000000 a0: 0x00000000000029a a1: 0x0000000000000001
a2: 0x0000000000000001 a3: 0x0000000000000002 a4: 0x0000000000000002 a5: 0x0000000000000007
a6: 0x0000000000000003 a7: 0x0000000000000001 s2: 0x0000000000000000 s3: 0x0000000000000000
s4: 0x0000000000000000 s5: 0x0000000000000000 s6: 0x0000000000000000 s7: 0x0000000000000000
s8: 0x0000000000000000 s9: 0x0000000000000000 s10: 0x0000000000000000 s11: 0x0000000000000000
t3: 0x00000000800039d0 t4: 0x0000000000000001 t5: 0x000000008000cdc1 t6: 0x0000000000000000
ft0: 0xffffffff00000000 ft1: 0xffffffff00000000 ft2: 0xffffffff00000000 ft3: 0xffffffff00000000
ft4: 0xffffffff00000000 ft5: 0xffffffff00000000 ft6: 0xffffffff00000000 ft7: 0xffffffff00000000
fs0: 0xffffffff00000000 fs1: 0xffffffff00000000 fa0: 0xffffffff00000000 fa1: 0xffffffff00000000
fa2: 0xffffffff00000000 fa3: 0xffffffff00000000 fa4: 0xffffffff00000000 fa5: 0xffffffff00000000
fa6: 0xffffffff00000000 fa7: 0xffffffff00000000 fs2: 0xffffffff00000000 fs3: 0xffffffff00000000
fs4: 0xffffffff00000000 fs5: 0xffffffff00000000 fs6: 0xffffffff00000000 fs7: 0xffffffff00000000
fs8: 0xffffffff00000000 fs9: 0xffffffff00000000 fs10: 0xffffffff00000000 fs11: 0xffffffff00000000
ft8: 0xffffffff00000000 ft9: 0xffffffff00000000 ft10: 0xffffffff00000000 ft11: 0xffffffff00000000
pc: 0x000000008000143a mstatus: 0x8000000a00006000 mcause: 0x0000000000000000 mepc: 0x00000004fe6a8c3dc
sstatus: 0x8000000200006000 scause: 0x0000000000000000 sepc: 0x0000000000000000
satp: 0x0000000000000000
mip: 0x0000000000000000 mie: 0x0000000000000000 mscratch: 0x0000000000000000 sscratch: 0x0000000000000000
mideleg: 0x0000000000000000 medeleg: 0x0000000000000000
mtval: 0x0000000000000000 stval: 0x9738662f59a3300f mtvec: 0x0000000000000000 stvec: 0x0000000000000000
privilege mode: 3 pmp: below
0: cfg:0x00 addr:0x0000000000000000| 1: cfg:0x00 addr:0x0000000000000000
2: cfg:0x00 addr:0x0000000000000000| 3: cfg:0x00 addr:0x0000000000000000
4: cfg:0x00 addr:0x0000000000000000| 5: cfg:0x00 addr:0x0000000000000000
6: cfg:0x00 addr:0x0000000000000000| 7: cfg:0x00 addr:0x0000000000000000
8: cfg:0x00 addr:0x0000000000000000| 9: cfg:0x00 addr:0x0000000000000000
10: cfg:0x00 addr:0x0000000000000000| 11: cfg:0x00 addr:0x0000000000000000
12: cfg:0x00 addr:0x0000000000000000| 13: cfg:0x00 addr:0x0000000000000000
14: cfg:0x00 addr:0x0000000000000000| 15: cfg:0x00 addr:0x0000000000000000
privilegeMode: 3
a3 different at pc = 0x008000180c, right= 0x0000000000000002, wrong = 0x0000000000000000
Core 0: ABORT at pc = 0x80000752
instrCnt = 1,547, cycleCnt = 12,103, IPC = 0.127820
Seed=0 Guest cycle spent: 12,106 (this will be different from cycleCnt if emu loads a snapshot)
Host time spent: 11,908ms
```

# Difftest

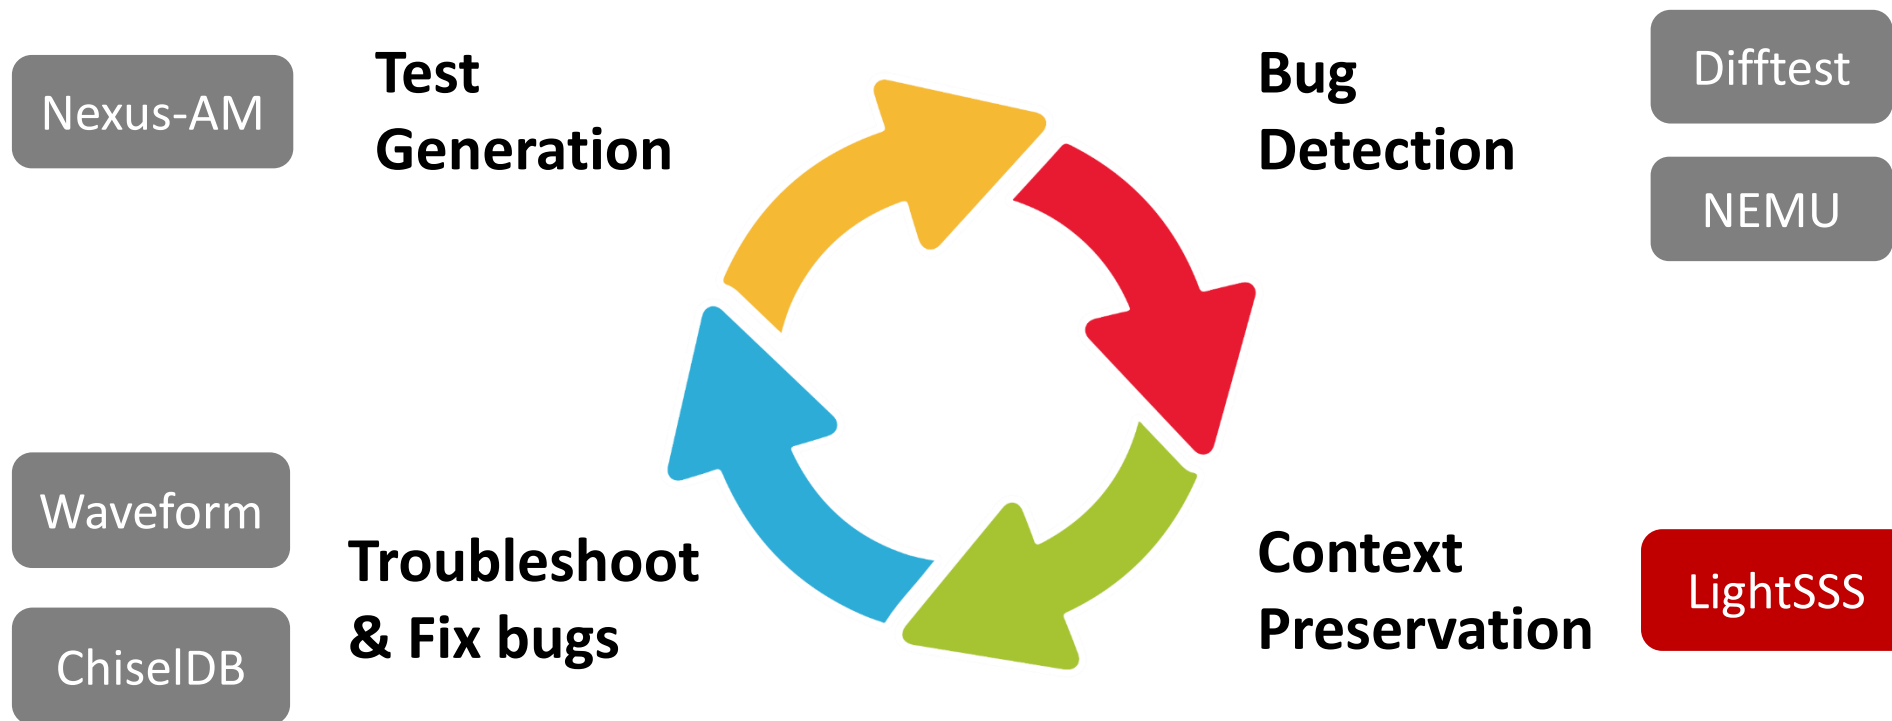
- **Hands-on:** Dump waveform after Difftest detects the bug
- **Showcase**

```
$ bash run-emu-dumpwave.sh
```

```
$ ls ../XiangShan/build | grep "vcd" # There should be a .vcd waveform
```

```
# ./emu-bug \  
-i $AM_HOME/apps/coremark/build/coremark-riscv64-xs.bin \  
--diff $NEMU_HOME/build/riscv64-nemu-interpretter-so \  
-b 12000 \      the begin cycle of the wave, you may find XiangShan trap at 12103s cycle in previous step  
-e 13000 \      the end cycle of the wave  
--dump-wave \   set this option to dump wave, you will find *.vcd on $NOOP_HOME/build
```


# MinJie Functional Verification





# Light-weight Simulation SnapShot: LightSSS

- Re-run simulation is time-consuming!
  - **Snapshot is the way out**
- LightSSS: snapshots of the process with fork()
  - Copy-on-write on Linux Kernel
- **Advantage 1: Good portability and scalability**
  - Snapshots for any external models (such as C++)
  - No need to understand details of external models
- **Advantage 2: Low overhead of snapshots**
  - ~500us for taking a snapshot
  - Far less overhead than RTL snapshots by Verilator





# LightSSS for on-demand debugging (waveform)

- **Hands-on:** use lightSSS to dump waveform
- **Showcase**

```
$ bash run-emu-lightsss.sh
```

```
# ./emu-bug \  
-i $AM_HOME/apps/coremark/build/coremark-riscv64-xs.bin \  
--diff $NEMU_HOME/build/riscv64-nemu-interpreter-so \  
--enable-fork \      No longer need to use complex parameters  
2 > lightsss.err
```



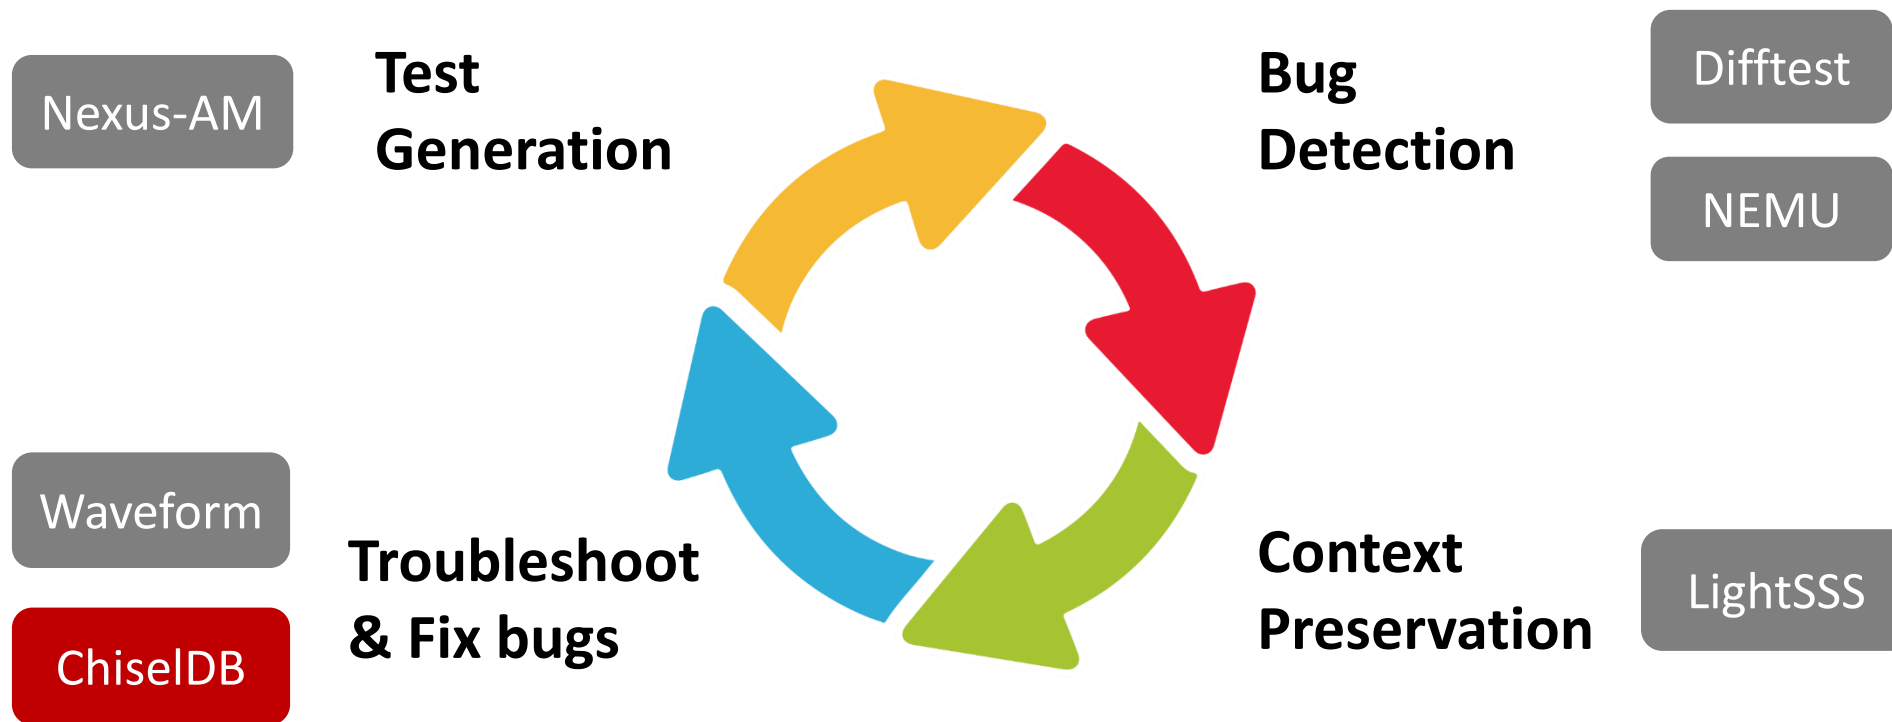
## LightSSS for on-demand debugging (waveform)

- **Hands-on:** use lightSSS to dump waveform
- LightSSS is working if you see:

```
[FORK_INFO pid(NUMBER)] the oldest checkpoint start to dump wave and dump nemu log...  
[FORK_INFO pid(NUMBER)] dump wave to /SOME/PATH/XXX.vcd...
```

- After that, simulation will start again from the existing latest snapshot

# MinJie Functional Verification







# Debug-friendly Structured Database: ChiselDB

- **Motivation:**

- Waveforms are **large** in size and **hard to apply** further analysis
- Need to analyze structured data like memory transaction trace

- **We present ChiselDB**

- **Highlights:**

- Inserting probes between module interfaces in hardware
- DPI-C: Using C++ function in Chisel code to transfer data
- Persist in database, SQL queries supported

- **Design source code:** *XiangShan/utility/src/main/scala/utility/ChiselDB.scala*
- **Usage:** Create table

```
// API: def createTable[T <: Record](tableName: String, hw: T): Table[T]
import huancun.utils.ChiselDB

class MyBundle extends Bundle {
    val fieldA = UInt(10.W)
    val fieldB = UInt(20.W)
}

val table = ChiselDB.createTable("MyTableName", new MyBundle)
```

- **Usage:** Register probes

```
/* APIs
def log(data: T, en: Bool, site: String = "", clock: Clock, reset: Reset)
def log(data: Valid[T], site: String, clock: Clock, reset: Reset): Unit
def log(data: DecoupledIO[T], site: String, clock: Clock, reset: Reset): Unit
*/
table.log(
  data = my_data /* hardware of type T */,
  en = my_cond,   site = "MyCallSite",
  clock = clock,  reset = reset
)
```

- **Example:** *XiangShan/src/main/scala/xiangshan/cache/mmu/L2TLB.scala*

```
class L2TlbMissQueueDB(implicit p: Parameters) extends TlbBundle {
  val vpn = UInt(vpnLen.W)
}

val L2TlbMissQueueInDB, L2TlbMissQueueOutDB = Wire(new L2TlbMissQueueDB)
L2TlbMissQueueInDB.vpn := missQueue.io.in.bits.vpn
L2TlbMissQueueOutDB.vpn := missQueue.io.out.bits.vpn

val L2TlbMissQueueTable = ChiselDB.createTable(
  "L2TlbMissQueue_hart" + p(XSCoreParamsKey).HartId.toString, new L2TlbMissQueueDB)

L2TlbMissQueueTable.log(L2TlbMissQueueInDB, missQueue.io.in.fire, "L2TlbMissQueueIn", clock, reset)
L2TlbMissQueueTable.log(L2TlbMissQueueOutDB, missQueue.io.out.fire, "L2TlbMissQueueOut", clock, reset)
```



- **Hands-on: Analysis of Cache Coherence Violation**
  - Add ChiselDB
  - Inject bug
  - Compile and run
  - Analyze ChiselDB

- Hands-on: Analysis of Cache Coherence Violation

*# Inject bug – We set all Release Data as a constant value*

```
$ cd chiseldb && cat cc_err.patch
```

```
--- a/src/main/scala/huancun/noninclusive/SinkC.scala
+++ b/src/main/scala/huancun/noninclusive/SinkC.scala
@@ -93,7 +93,7 @@ class SinkC(implicit p: Parameters) extends
BaseSinkC {
-    buffer(insertIdx)(count) := c.bits.data
+    buffer(insertIdx)(count) := 0xABCDEF.U
```

- Hands-on: Analysis of Cache Coherence Violation

```
# simulate using pre-built emu (ChiselDB enabled)
```

```
$ bash step1-run.sh
```

```
# ./emu-cc-err -i $NOOP_HOME/ready-to-run/microbench.bin \  
--diff $NOOP_HOME/ready-to-run/riscv64-nemu-interpreter-so \  
--dump-db # set this argument to dump data base  
2>linux.err
```

- Hands-on: Analysis of Cache Coherence Violation

```
# Analyze
```

```
$ bash step2-analyze.sh
```

```
# 1. sqlite query for all transactions on Eaddr
```

```
# 2. use script to parse TLLog
```

```
# sqlite3 $NOOP_HOME/build/*.db \
```

```
  "select * from TLLOG where ADDRESS=0x80008fc0" | \
```

```
sh $NOOP_HOME/scripts/utls/parseTLLog.sh
```





# ChiselDB:

*# Result: [Time | To\_From | Channel | Opcode | Permission | Address | Data(0)]*

*# Data successfully transferred from L1D to L2*

*186297 L2\_L1D\_0 C ReleaseData Shrink TtoN 80008fc0 0000000000000000a*


*# But when L1D acquires addr again, data loaded from L2 is wrong*

*187022 L2\_L1D\_0 A AcquireBlock Grow NtoB 80008fc0*

*187029 L2\_L1D\_0 D GrantData Cap toT 80008fc0 0000000000abcdef*

*# So there must be something wrong when L2 records Release Data*


# MinJie Performance Verification





# Checkpoints

- Uniform Checkpoint
- Divide a program to small parts which can be simulated in parallel



- Simpoint Checkpoint
- Select representative parts from a program with cluster analysis



# Checkpoints: Uniform Checkpoint

- step0: Prepare NEMU environment for checkpoints

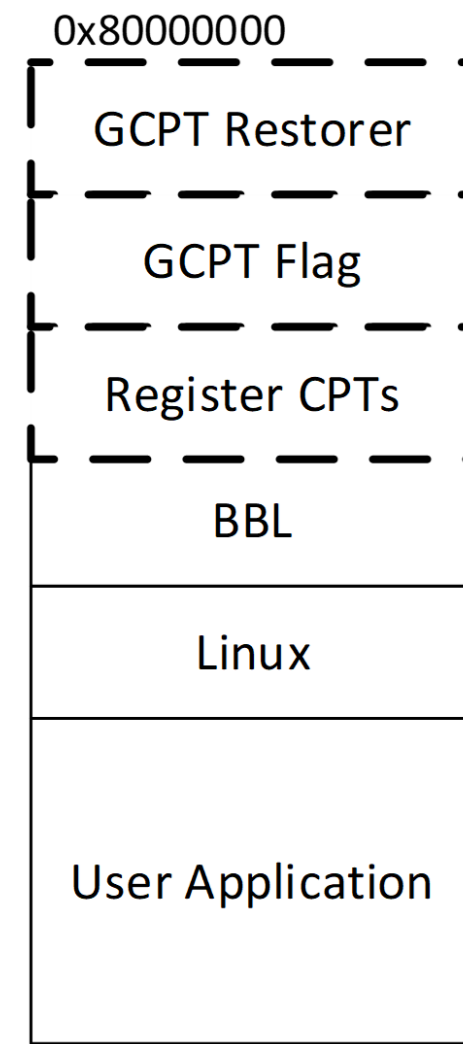
```
$ cd ../checkpoint          # tutorial/checkpoint
$ bash uniform-step0-prepare.sh
```

```
# uniform-step0-prepare.sh

# cd $NEMU_HOME

# git checkout asplos2023-tutorial
# git submodule update --init
# make clean
# make tutorial_defconfig
# make -j4

// generate gcpt restorer binary
# cd resource/gcpt_restore && make -j4
```



Basic format of RISC-V Checkpoint

# Checkpoints: Uniform Checkpoint

- step1: Generate uniform checkpoints in NEMU

```
$ bash uniform-step1-gen.sh
```

```
# uniform-step1-gen.sh

# cd $NEMU_HOME
# rm -rf tutorial_uniform
# ./build/riscv64-nemu-interpretor
#   --cpt-interval 2000000  the profiling period for simpoint
#   -u                    uniformly take cpt with fixed interval
#   -b                    run with batch mode
#   -D tutorial_uniform    directory where the checkpoint is generated
#   -C try                 name of this task
#   -w linux               name of workload
#   -r ./resource/gcpt_restore/build/gcpt.bin  binary of gcpt restorer
#   --dont-skip-boot       checkpoint immediately after boot
#   -I 3000000             max number of instructions executed
#   ./ready-to-run/linux-0xa0000.bin
```

# 

```
[ 0.000000] Dentry cache hash table entries: 4096 (order: 3, 32768 bytes)
[ 0.000000] Inode cache hash table entries: 2048 (order: 2, 16384 bytes)
[ 0.000000] Sorting __ex_table...
src/checkpoint/serializer.cpp:255,shouldTakeCpt] Should take cpt now: 2000002
src/checkpoint/path_manager.cpp:77,setOutputDir] Created tutorial_uniform/try/linux/0/

[src/checkpoint/serializer.cpp:127,serializeRegs] Writing int registers to checkpoint memory @[0x80001000, 0x80001100) [0x1000, 0x1100)
[src/checkpoint/serializer.cpp:137,serializeRegs] Writing float registers to checkpoint memory @[0x80001100, 0x80001200) [0x1100, 0x1200)
[src/checkpoint/serializer.cpp:145,serializeRegs] Writing PC: 0xffffffff80003e4c at addr 0x80001200
[src/checkpoint/serializer.cpp:168,serializeRegs] CSR 0x105: 0xffffffff8039075c
[src/checkpoint/serializer.cpp:168,serializeRegs] CSR 0x106: 0xffffffffffffffff
[src/checkpoint/serializer.cpp:168,serializeRegs] CSR 0x141: 0x80200098
[src/checkpoint/serializer.cpp:168,serializeRegs] CSR 0x142: 0xc
[src/checkpoint/serializer.cpp:168,serializeRegs] CSR 0x143: 0x80200098
[src/checkpoint/serializer.cpp:168,serializeRegs] CSR 0x180: 0x800000000000806f
[src/checkpoint/serializer.cpp:168,serializeRegs] CSR 0x300: 0xa00000180
[src/checkpoint/serializer.cpp:168,serializeRegs] CSR 0x301: 0x8000000000014112d
[src/checkpoint/serializer.cpp:168,serializeRegs] CSR 0x302: 0xb109
[src/checkpoint/serializer.cpp:168,serializeRegs] CSR 0x303: 0x222
[src/checkpoint/serializer.cpp:168,serializeRegs] CSR 0x304: 0x22a
[src/checkpoint/serializer.cpp:168,serializeRegs] CSR 0x305: 0x800a0004
[src/checkpoint/serializer.cpp:168,serializeRegs] CSR 0x306: 0xffffffffffffff
[src/checkpoint/serializer.cpp:168,serializeRegs] CSR 0x340: 0x800abdc0
[src/checkpoint/serializer.cpp:168,serializeRegs] CSR 0x341: 0xffffffff80390e4d
[src/checkpoint/serializer.cpp:168,serializeRegs] CSR 0x342: 0x9
[src/checkpoint/serializer.cpp:168,serializeRegs] CSR 0x3a0: 0x1f0800
[src/checkpoint/serializer.cpp:168,serializeRegs] CSR 0x3b0: 0x20000000
[src/checkpoint/serializer.cpp:168,serializeRegs] CSR 0x3b1: 0x2002d000
[src/checkpoint/serializer.cpp:168,serializeRegs] CSR 0x3b2: 0x3fffffff
[src/checkpoint/serializer.cpp:168,serializeRegs] CSR 0x5c4: 0x3
[src/checkpoint/serializer.cpp:171,serializeRegs] Writing CSR to checkpoint memory @[0x80001300, 0x800019300) [0x1300, 0x9300)
[src/checkpoint/serializer.cpp:179,serializeRegs] Touching flag: 0xbccf at addr 0x80000f00
[src/checkpoint/serializer.cpp:183,serializeRegs] Record mode flag: 0x1 at addr 0x80000f08
[src/checkpoint/serializer.cpp:188,serializeRegs] Record time: 0x1 at addr 0x80000f10
[src/checkpoint/serializer.cpp:192,serializeRegs] Record time: 0x1 at addr 0x80000f18
[src/checkpoint/serializer.cpp:81,serializePMem] Put gcpt restorer ./resource/gcpt_restore/build/gcpt.bin to start of pmem
Opening tutorial_uniform/try/linux/0/_2000002_.gz as checkpoint output file
[src/checkpoint/serializer.cpp:110,serializePMem] Written 0x7ffffff bytes
[src/checkpoint/serializer.cpp:110,serializePMem] Written 0x7ffffff bytes
[src/checkpoint/serializer.cpp:110,serializePMem] Written 0x7ffffff bytes
[src/checkpoint/serializer.cpp:110,serializePMem] Written 0x4 bytes
[src/checkpoint/serializer.cpp:116,serializePMem] Checkpoint done!

[src/checkpoint/serializer.cpp:263,notify_taken] Taking checkpoint @ instruction count 2000002
[src/cpu/cpu-exec.c:203,per_bb_profile] Should take checkpoint on pc 0xffffffff80003e4c
[ 0.000000] Memory: 25548K/30720K available (699K kernel code, 78K rdata, 102K rodata, 3648K init, 98K bss, 5172K reserved, 0K cma-reserved)
[ 0.000000] SLUB: Hwalign=64, Order=0-3, MinObjects=0, CPUs=1, Nodes=1
[ 0.000000] NR_IRQS: 0, nr_irqs: 0, preallocated irq: 0
```

Take a checkpoint from here

Save int and float registers

Save CSR registers

Save other information and gcpt restorer

Write checkpoints to gz file

# Checkpoints: Uniform Checkpoint

- **step2: Run the uniform checkpoint in NEMU or XiangShan**
- Here we take NEMU as an example

```
$ bash uniform-step2-run-nemu.sh
```

```
# uniform-step2-run-nemu.sh
```

```
# cd $NEMU_HOME
```

```
# ./build/riscv64-nemu-interpret
```

```
# -b run with batch mode
```

```
# --restore restoring from CPT FILE
```

```
# -I 1000000 max number of instructions executed
```

```
# `find tutorial_uniform/try/linux/0/ -type f -name "*.gz" | tail -1`  
path of the checkpoint files just generated
```

# Checkpoints: Uniform Checkpoint

```
[src/monitor/monitor.c:103,parse_args] Restoring from checkpoint
[src/device/io/mmio.c:38,add_mmio_map] Add mmio map 'clint' at [0x0000000003800000, 0x0000000003800fff]
[src/isa/riscv64/init.c:100,init_isa] NEMU will start from pc 0x80000000
[src/monitor/image_loader.c:69,load_img] Loading Gcpt file form cmdline: tutorial_uniform/try/linux/0/_2000002_.gz

[src/monitor/image_loader.c:77,load_img] Loading GZ image tutorial_uniform/try/linux/0/_2000002_.gz
[src/monitor/image_loader.c:69,load_img] Loading Gcpt restorer form cmdline: (null)

[src/monitor/image_loader.c:71,load_img] No image is given. Use the default built-in image/restorer.
[src/device/io/port-io.c:30,add_pio_map] Add port-io map 'uartlite' at [0x000000000000003f8, 0x0000000000000404]
[src/device/io/mmio.c:38,add_mmio_map] Add mmio map 'uartlite' at [0x0000000004060000, 0x000000000406000c]
[src/device/io/port-io.c:30,add_pio_map] Add port-io map 'rtc' at [0x0000000000000048, 0x000000000000004f]
[src/device/io/mmio.c:38,add_mmio_map] Add mmio map 'rtc' at [0x000000000a100048, 0x000000000a10004f]
[src/device/io/port-io.c:30,add_pio_map] Add port-io map 'screen' at [0x0000000000000100, 0x0000000000000107]
[src/device/io/mmio.c:38,add_mmio_map] Add mmio map 'screen' at [0x0000000004000100, 0x0000000004000107]
[src/device/io/mmio.c:38,add_mmio_map] Add mmio map 'vmem' at [0x0000000005000000, 0x000000000500752ff]
[src/device/io/port-io.c:30,add_pio_map] Add port-io map 'keyboard' at [0x0000000000000060, 0x0000000000000063]
[src/device/io/mmio.c:38,add_mmio_map] Add mmio map 'keyboard' at [0x000000000a100060, 0x000000000a100063]
[src/device/io/mmio.c:38,add_mmio_map] Add mmio map 'sdhci' at [0x0000000004000200, 0x0000000004000207f]
[src/device/sdcard.c:137,init_sdcard] Can not find sdcard image:
[src/monitor/monitor.c:44,welcome] Debug: OFF
[src/monitor/monitor.c:49,welcome] Build time: 19:35:10, Mar 25 2023
Welcome to riscv64-NEMU!
For help, type "help"
[ 0.000000] Memory: 25548K/30720K available (609K kernel code, 78K rwd, 102K rodata, 3648K init, 98K bss, 5172K reserved, 0K cma-reserved)
[ 0.000000] SLUB: HWalign=64, Order=0-3, MinObjects=0, CPUs=1, Nodes=1
[ 0.000000] NR_IRQS: 0, nr_irqs: 0, preallocated irq: 0
[ 0.000000] clocksource: riscv_clocksource: mask: 0xffffffffffffffff max_cycles: 0x1d854df40, max_idle_ns: 3526361616960 ns
[ 0.000000] console [hvc0] enabled
[ 0.000000] console [hvc0] enabled
[ 0.000000] bootconsole [early0] disabled
[ 0.000000] bootconsole [early0] disabled
[ 0.000000] Calibrating delay loop (skipped), value calculated using timer frequency.. 2.00 BogoMIPS (lpj=10000)
[ 0.000000] pid_max: default: 4096 minimum: 301
[ 0.000000] Mount-cache hash table entries: 512 (order: 0, 4096 bytes)
[ 0.000000] Mountpoint-cache hash table entries: 512 (order: 0, 4096 bytes)
[ 0.010000] clocksource: jiffies: mask: 0xffffffff max_cycles: 0xffffffff, max_idle_ns: 19112604462750000 ns
[ 0.020000] clocksource: Switched to clocksource riscv_clocksource
```

Restore from checkpoint  
and load gz files

NEMU will start execution from  
the moment of that checkpoint



# Checkpoints: Simpoint Checkpoint

- step0: Prepare NEMU environment for Simpoint checkpoint

```
$ bash simpoint-step0-prepare.sh
```

```
# simpoint-step0-prepare.sh
```

```
# cd $NEMU_HOME
```

```
# git checkout cpt-bk
```

```
# git submodule update --init
```

```
# cd resource/simpoint/simpoint_repo/analysiscode && make simpoint -j4
```

*generate simpoint generator binary*

```
# cd resource/gcpt_restore && make -j4
```

*generate gcpt restorer binary*

```
# cd $NEMU_HOME
```

```
# make clean
```

```
# make ISA=riscv64 XIANGSHAN=1 -j4
```

*compile NEMU*

# Checkpoints: Simpoint Checkpoint

- **step1: Execute the workload and collect program behavior**

```
$ bash simpoint-step1-profiling.sh
```

```
# simpoint-step1-profiling.sh
```

```
# cd $NEMU_HOME
```

```
# rm -rf tutorial_simpoint
```

```
# ./build/riscv64-nemu-interpret
```

```
# -b run with batch mode
```

```
# -D $NEMU_HOME/tutorial_simpoint directory where the checkpoint is generated
```

```
# -C profiling name of this task
```

```
# -w stream name of workload
```

```
# --simpoint-profile is simpoint profiling
```

```
# --interval 50000000 simpoint interval instructions
```

```
# $XS_PROJECT_ROOT/tutorial/bin/stream-0xa0000.bin
```

# Checkpoints: Simpoint Checkpoint

```
[src/isa/riscv64/exec/special.c,38,exec_nemu_trap] Start profiling, resetting inst count
```

```
-----  
STREAM version $Revision: 5.10 $  
-----
```

```
This system uses 4 bytes per array element.  
-----
```

```
Array size = 2000000 (elements), Offset = 0 (elements)
```

```
Memory per array = 7.6 MiB (= 0.0 GiB).
```

```
Total memory required = 22.9 MiB (= 0.0 GiB).
```

```
Each kernel will be executed 10 times.
```

```
The *best* time for each kernel (excluding the first iteration)  
will be used to compute the reported bandwidth.  
-----
```

```
Your clock granularity/precision appears to be 2047 microseconds.
```

```
Each test below will take on the order of 102400 microseconds.  
(= 50 clock ticks)
```

```
Increase the size of the arrays if this shows that  
you are not getting at least 20 clock ticks per test.  
-----
```

```
WARNING -- The above is only a rough guideline.
```

```
For best results, please be sure you know the  
precision of your system timer.  
-----
```

Function	Best	Rate MB/s	Avg time	Min time	Max time
Copy:	710.2		0.023438	0.022528	0.024576
Scale:	181.7		0.089884	0.088064	0.092160
Add:	217.0		0.112868	0.110592	0.114688
Triad:	172.3		0.140402	0.139264	0.141312

```
-----
```

```
Solution Validates: avg error less than 1.000000e-06 on all three arrays  
-----
```

```
[src/monitor/cpu-exec.c,110,cpu_exec] nemu: HIT GOOD TRAP at pc = 0x0000002aaaaaa5f6
```

NEMU will do profiling while  
executing workload STREAM

# Checkpoints: Simpoint Checkpoint

- step2: Cluster and obtain multiple representative slices

```
$ bash simpoint-step2-cluster.sh
```

```
# simpoint-step2-cluster.sh
```

```
# cd $NEMU_HOME
```

```
# mkdir -p $NEMU_HOME/tutorial_simpoint/cluster/stream
```

```
# export CLUSTER=$NEMU_HOME/tutorial_simpoint/cluster/stream
```

```
# ./resource/simpoint/simpoint_repo/bin/simpoint
```

```
# -loadFVFile ./tutorial_simpoint/profiling/simpoint_bbv.gz load a frequency vector file
```

```
# -saveSimpoints $CLUSTER/simpoints0 file to save simpoints
```

```
# -saveSimpointWeights $CLUSTER/weights0 file to save simpoint weights
```

```
# -inputVectorsGzipped input vectors have been compressed with gzip
```

```
# -maxK 3 maximum number of clusters to use
```

```
# -numInitSeeds 2 times of different random initialization for each run, taking only the best clustering
```

```
# -iters 1000 maximum number of iterations that should perform
```

```
# -seedkm 123456 random seed for choosing initial k-means centers
```

```
# -seedproj 654321 random seed for random linear projection
```

# Checkpoints: Simpoint Checkpoint

```
-----  
Run number 3 of at most 4, k = 2  
-----
```

```
-----  
Initialization seed trial #1 of 2; initialization seed = 123460  
-----
```

```
-----  
Initialized k-means centers using random sampling: 2 centers  
-----
```

```
Number of k-means iterations performed: 1
```

```
BIC score: 123.793
```

```
Distortion: 1.61246
```

```
Distortions/cluster: 1.54841 0.0640483
```

```
Variance: 0.161246
```

```
Variances/cluster: 1.54841 0.00711647  
-----
```

Compute the information  
of K-means clustering

```
-----  
Initialization seed trial #2 of 2; initialization seed = 123461  
-----
```

```
-----  
Initialized k-means centers using random sampling: 2 centers  
-----
```

```
Number of k-means iterations performed: 1
```

```
BIC score: 123.793
```

```
Distortion: 1.61246
```

```
Distortions/cluster: 0.0640483 1.54841
```

```
Variance: 0.161246
```

```
Variances/cluster: 0.00711647 1.54841  
-----
```

Obtain multiple program  
representative checkpoints

```
The best initialization seed trial was #1  
-----
```

```
-----  
Post-processing runs  
-----
```

```
-----  
For the BIC threshold, the best clustering was run 3 (k = 2)  
-----
```

```
Post-processing run 3 (k = 2)  
-----
```

# Checkpoints: Simpoint Checkpoint

- **step3: Generate corresponding Checkpoints based on clustering results**

```
$ bash simpoint-step3-take-cpt.sh          # It may take about 3 minutes
```

```
# simpoint-step3-take-cpt.sh

# cd $NEMU_HOME
# rm -rf $NEMU_HOME/tutorial_simpoint/stream/take_cpt

# ./build/riscv64-nemu-interpretter
#   -b                               run with batch mode
#   -D tutorial_simpoint             directory where the checkpoint is generated
#   -C take_cpt                     name of this task
#   -w stream                        name of workload
#   -S ./tutorial_simpoint/cluster  simpoints dir
#   --checkpoint-interval 50000000  simpoint interval instructions
#   $XS_PROJECT_ROOT/tutorial/bin/stream-0xa0000.bin
```



# Checkpoint: Simpoint Checkpoint

```
[src/isa/riscv64/exec/special.c,38,exec_nemu_trap] Start profiling, resetting inst count
[src/checkpoint/serializer.cpp,100,serializeRegs] Writing int registers to checkpoint memory @[0x80001000, 0x80001100) [0x1000, 0x1100)
[src/checkpoint/serializer.cpp,110,serializeRegs] Writing float registers to checkpoint memory @[0x80001100, 0x80001200) [0x1100, 0x1200)
[src/checkpoint/serializer.cpp,118,serializeRegs] Writing PC: 0x7fffffe000720720 at addr 0x80001200
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x105: 0xffffffff000697cc8
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x106: 0xffffffffffffffff
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x141: 0x200013417e
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x142: 0x8
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x180: 0x800000000000f805d
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x300: 0xa00000022
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x301: 0x800000000014112d
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x302: 0xb109
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x303: 0x222
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x304: 0x22a
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x305: 0x800a0004
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x306: 0xffffffffffffffff
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x340: 0x800abdc0
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x341: 0xffffffff0008d6eb2
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x342: 0x2
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x3a0: 0x1f0800
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x3b0: 0x20028000
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x3b1: 0x2002d000
[src/checkpoint/serializer.cpp,141,serializeRegs] CSR 0x3b2: 0xffffffffffffffff
[src/checkpoint/serializer.cpp,144,serializeRegs] Writing CSR to checkpoint memory @[0x80001300, 0x80001300) [0x1300, 0x1300)
[src/checkpoint/serializer.cpp,152,serializeRegs] Touching Flag: 0xb0ef at addr 0x80001f00
[src/checkpoint/serializer.cpp,156,serializeRegs] Record mode flag: 0x1 at addr 0x80001f08
[src/checkpoint/serializer.cpp,160,serializeRegs] Record time: 0x1 at addr 0x80001f10
[src/checkpoint/serializer.cpp,164,serializeRegs] Record time: 0x1 at addr 0x80001f18
[src/checkpoint/serializer.cpp,52,serializePMem] Created tutorial_simpoint/take_cpt/stream_0_0.083333/0/

Opening tutorial_simpoint/take_cpt/stream_0_0.083333/0/_0_.gz as checkpoint output file
[src/checkpoint/serializer.cpp,83,serializePMem] Written 0x7fffffff bytes
[src/checkpoint/serializer.cpp,83,serializePMem] Written 0x7fffffff bytes
[src/checkpoint/serializer.cpp,83,serializePMem] Written 0x7fffffff bytes
[src/checkpoint/serializer.cpp,83,serializePMem] Written 0x4 bytes
[src/checkpoint/serializer.cpp,89,serializePMem] Checkpoint done!
```

Start profiling, resetting inst count → Start making checkpoint

Save int & float registers

Save CSR registers

Save other information

write checkpoints into .gz

## Checkpoints: Simpoint Checkpoint

- **step4: Run the simpoint checkpoint in NEMU or XiangShan**
- Here we take XiangShan as an example

```
$ bash simpoint-step4-run-xs.sh
```

```
# simpoint-step4-run-xs.sh
```

```
$XS_PROJECT_ROOT/tutorial/emu  
-i `find $NEMU_HOME/tutorial_simpoint/ -type f -name "*.gz" | tail -1`  
--diff $NOOP_HOME/ready-to-run/riscv64-nemu-interpreter-so  
--max-cycles=100000  
2>simpoint.err
```



# Simulation Performance Counter

- **Two common types:**
  - **Accumulation**: basic accumulator counter
  - **Histogram**: count the distribution

## Accumulation:

```
[PERF][time=10000] ctrlBlock.rob: commitInstr,      1500  
[PERF][time=10000] ctrlBlock.rob: waitLoadCycle,    800
```

## Histogram:

```
[PERF][time=10000] l2cache: acquire_period_10_20,   15  
[PERF][time=10000] l2cache: acquire_period_20_30,  200  
[PERF][time=10000] l2cache: acquire_period_30_40,   60
```

# Simulation Performance Counter

- Usage:

- Add `XSPerfAccumulate('name', signal)` or *Histogram* wherever you want
- **By default**, it is printed after simulation finishes
- Set `DebugOptions(EnablePerfDebug = false)` to turn off

- Showcase

```
$ cd ..
```

```
$ bash run-emu-perf.sh
```

```
# ./emu -i hello.bin \
```

```
--diff $NEMU_HOME/build/riscv64-nemu-interpreter-so 2>perf.err
```

```
$ less perf.err
```

# Simulation Performance Counter

```
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.ctrlBlock.decode: utilization,          1183
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.bpu.predictors.components_3.tables_0.wrbypass: wrbypass_hit,          0
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.icache.missUnit.entries_1: entryPenalty1,          714
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.ctrlBlock.decode: waitInstr,          414
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.bpu.predictors.components_3.tables_0.wrbypass: wrbypass_miss,          0
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.ctrlBlock.decode: stall_cycle,          393
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.icache.missUnit.entries_1: entryReq1,          14
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.bpu.predictors.components_3.tables_1.wrbypass: wrbypass_hit,          0
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.frontend: FrontendBubble,          4085
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.icache.missUnit.entries_0: entryPenalty0,          248
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.bpu.predictors.components_3.tables_1.wrbypass: wrbypass_miss,          2
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.ibuffer: utilization,          1826
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.icache.missUnit.entries_0: entryReq0,          6
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.ibuffer: util_0_1,          2029
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.ibuffer: util_1_2,          541
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.misc.busPMU: dcache_bank_0_A_channel_fire,          8
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.ibuffer: util_2_3,          45
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.misc.busPMU: dcache_bank_0_A_channel_stall,          0
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.ibuffer: util_3_4,          107
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.exuBlocks.scheduler.rs_4.staRS_0.oldestSelection: oldest_override_0,          0
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.misc.busPMU: dcache_bank_0_A_channel_PutFullData_fire,          0
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.ibuffer: util_4_5,          44
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.exuBlocks.scheduler.rs_4.staRS_0.oldestSelection: oldest_same_as_selected_0,          0
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.misc.busPMU: dcache_bank_0_A_channel_PutFullData_stall,          0
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.ibuffer: util_5_6,          85
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.exuBlocks.scheduler.rs_4.staRS_0.oldestSelection: oldest_override_1,          0
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.misc.busPMU: dcache_bank_0_A_channel_PutPartialData_fire,          0
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.ibuffer: util_6_7,          13
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.exuBlocks.scheduler.rs_4.staRS_0.oldestSelection: oldest_same_as_selected_1,          0
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.misc.busPMU: dcache_bank_0_A_channel_PutPartialData_stall,          0
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.ibuffer: util_7_8,          27
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.memBlock.dtlb_st_tlb_st.entries: port0_np_sp_multi_hit,          0
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.misc.busPMU: dcache_bank_0_A_channel_ArithmeticData_fire,          0
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.ibuffer: full,          862
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.memBlock.dtlb_st_tlb_st.entries: port1_np_sp_multi_hit,          0
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.misc.busPMU: dcache_bank_0_A_channel_ArithmeticData_stall,          0
[PERF ][time=      2893] TOP.SimTop.l_soc.core_with_l2.core.frontend.ibuffer: exHalf,          125
```

Find more parsing scripts at <https://github.com/OpenXiangShan/env-scripts/blob/main/perf/perf.py>

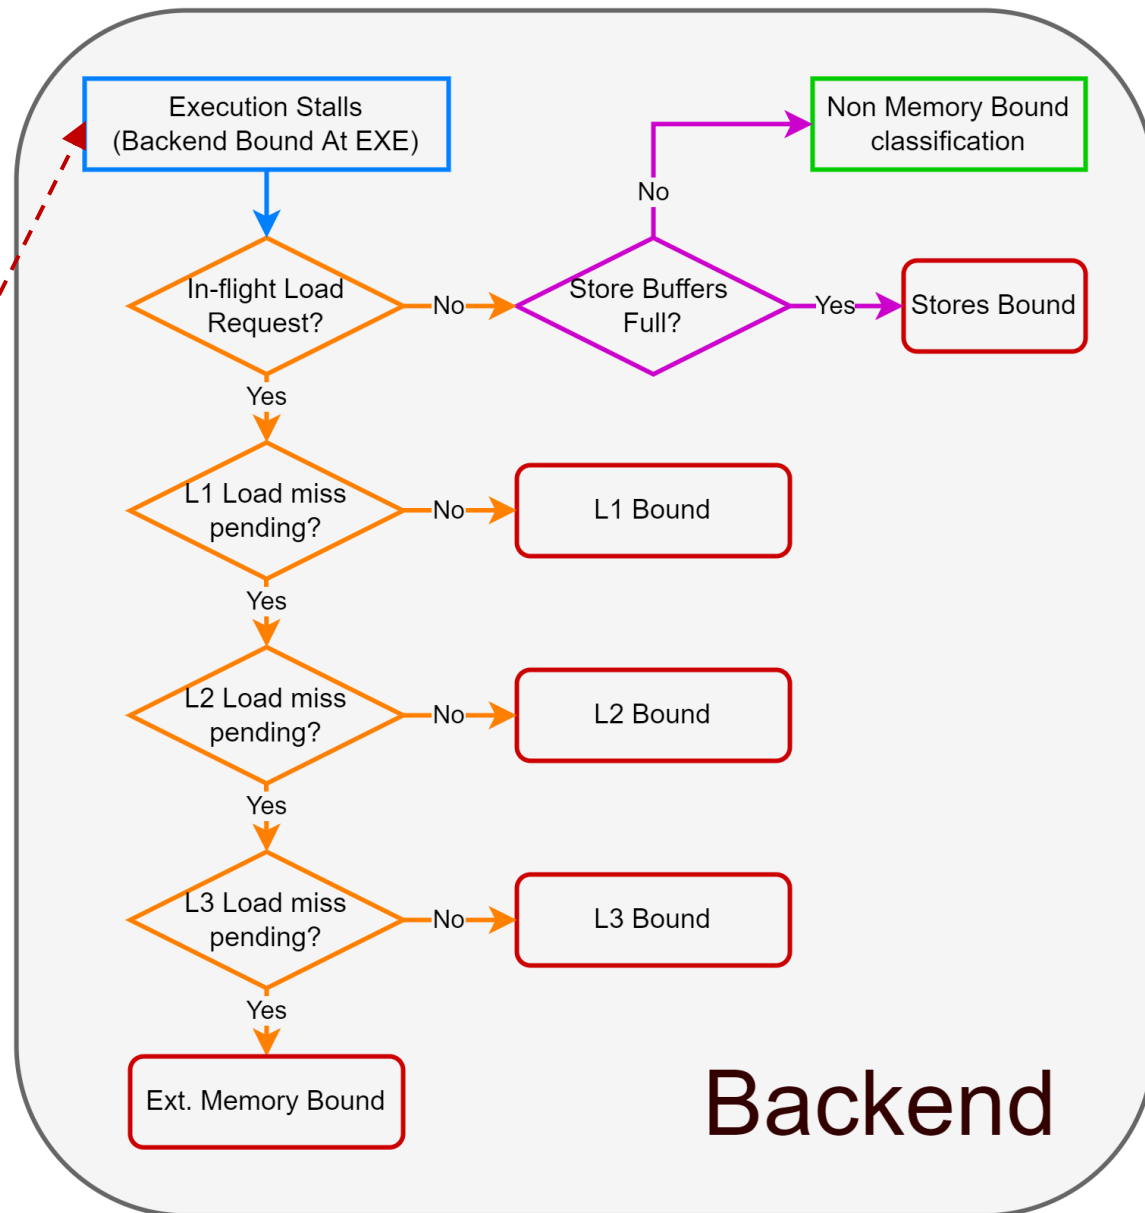
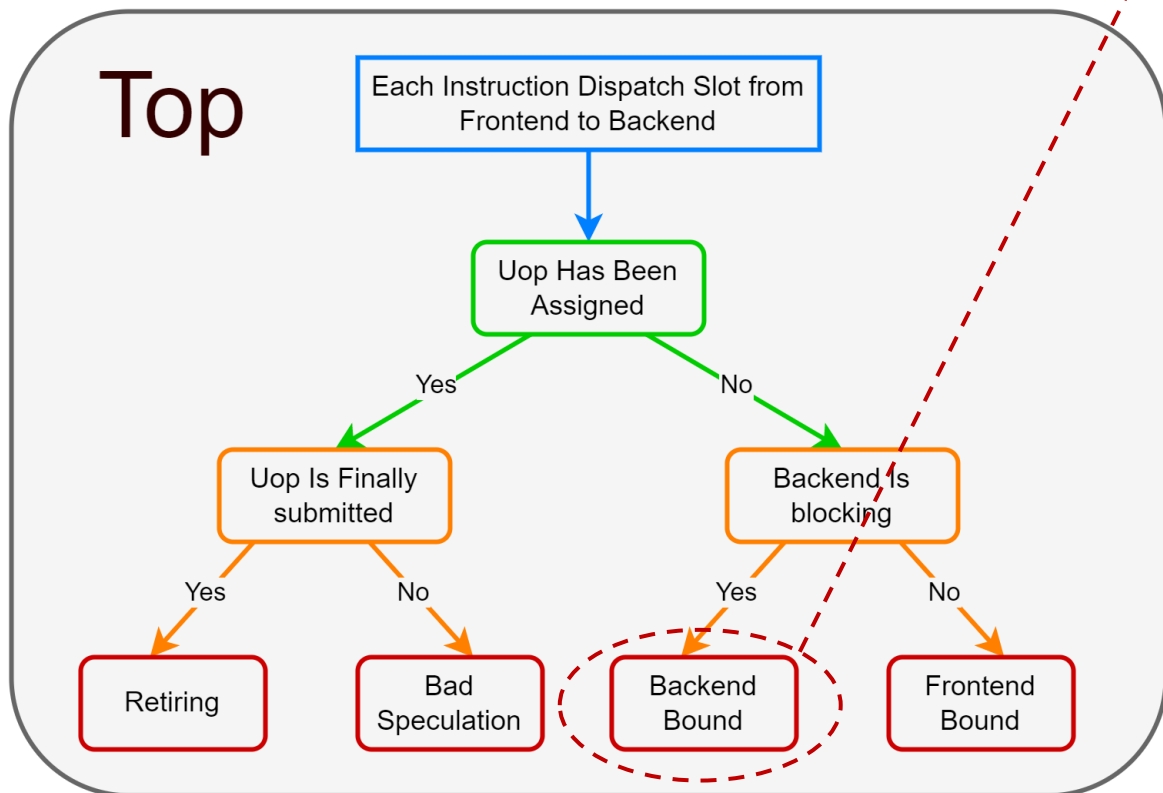


# Top-Down: Method for Performance Analysis

- Organize scattered performance events in a **hierarchical manner**
- **Accurately** calculate one event's impact on processor performance
- We apply the Top-Down approach to XiangShan
  - Complete **targeted optimization adaptation** for RISC-V instruction set
  - Optimize **performance counters** according to XiangShan microarchitecture
  - Further refine the **hierarchical design** of Top-Down model
  - Without missing or duplicating any events
  - Without assuming the blocking cycles of performance events in advance
- Due to the limited time, we will not arrange a hands-on section here
  - Please refer to <https://github.com/OpenXiangShan/XiangShan/tree/master/scripts/top-down>

# Top-Down

- Use Top and Backend as examples





# Top-Down

- Setting Performance Counters in RTL Code

*Dispatch.scala*

```
val stallReason = Wire(chiselTypeOf(io.stallReason.reason))
// ...
TopDownCounters.values.foreach(ctr =>
  XSPerfAccumulate(ctr.toString(), PopCount(stallReason.map(_ === ctr.id.U)))
)
```

- Do simulation and collect performance counter data


# Top-Down

- Analyze through Top-Down method

```
env-scripts: perf/top_down/configs.py
```

```
xs_coarse_rename_map = {  
    'OverrideBubble': 'MergeFrontend',  
    'FtqFullStall': 'MergeFrontend',  
    'IntDqStall': 'MergeCoreDQStall', # attribute perf-counters to Top-Down hierarchy  
    'FpDqStall': 'MergeCoreDQStall'  
}
```

- Obtain analysis results and make targeted optimizations







# ConstantIn

- **Motivation**

- Wanna test the performance under different parameters
- But re-compilation is **time consuming**
- Can we change parameters(constants) at runtime?

- **We present Constantin** (Constant - in)

- **Design/Highlights:**

- DPI-C: Using C++ function in Chisel code (through Verilog using BlackBox)
- Signal values configured at runtime rather than compile time



- Constants keep constant only during runtime

```
+ XiangShan_another git:(constantinople) * cat build/constantin.txt
lowerbound 10
predSel 1

+ XiangShan_another git:(constantinople) * ./build/emu -i ~/oracle/XiangShan_oracle/ready-to-run/microbench.bin -I 10000 -e 0 2> /dev/null
Emu compiled at Feb 6 2023, 11:24:27
Using simulated 327688 Flash
No valid flash bin path, use preset flash instead
The image is /nfs/home/chenguokai/oracle/XiangShan_oracle/ready-to-run/microbench.bin
Using simulated 8192MB RAM
--diff is not given, try to use $(NEMU_HOME)/build/riscv64-nemu-interpretor-so by default
NemuProxy using /nfs/home/chenguokai/NEMU_backup/build/riscv64-nemu-interpretor-so
The first instruction of core 0 has committed. DiffTest enabled.
[NEMU] Can not find flash
[NEMU] Use built-in image
[soc/device/io/mmio.c:38]
===== Running Microbench
[qsart] Quick sort: Core
total guest instructions
instrCnt = 50,001, cycleCnt
Seed0 Guest cycle spent: 0.000000
Host time spent: 0.000000
+ XiangShan_another git:
lowerbound 10
predSel 1

+ XiangShan_another git:
bin -I 10000 -e 0 2> /dev/null
Emu compiled at Feb 6 2023
Using simulated 327688 Flash
No valid flash bin path, use preset flash instead
The image is /nfs/home/chenguokai/oracle/XiangShan_oracle/ready-to-run/microbench.bin
Using simulated 8192MB RAM
--diff is not given, try to use $(NEMU_HOME)/build/riscv64-nemu-interpretor-so by default
NemuProxy using /nfs/home/chenguokai/NEMU_backup/build/riscv64-nemu-interpretor-so
Assertion failed at line 722859
The simulation stopped. There might be some assertion failed.
Core 0: abort at pc = 0xfffff937b1e4
total guest instructions = 0
instrCnt = 0, cycleCnt = 3, IPC = 0.000000
Seed0 Guest cycle spent: 6 (this will be different from cycleCnt if emu loads a snapshot)
Host time spent: 4ms
```

Debugging

```
+ XiangShan git:(master) * export NEMU_HOME=/nfs/home/chenguokai/NEMU_backup
+ XiangShan git:(master) * export MDP_HOME=/nfs/home/chenguokai/tutorial_demo/Constantin/XiangShan
+ XiangShan git:(master) * python3 scripts/constantin.py constant.json
iteration 0 begins

[[['block_cycles_cache_0', 6], ['block_cycles_cache_1', 42], ['block_cycles_cache_2', 19], ['block_cycles_cache_3', 5]]
[[['block_cycles_cache_0', 12], ['block_cycles_cache_1', 48], ['block_cycles_cache_2', 29], ['block_cycles_cache_3', 4]]
[[['block_cycles_cache_0', 13], ['block_cycles_cache_1', 19], ['block_cycles_cache_2', 31], ['block_cycles_cache_3', 2]]
[[['block_cycles_cache_0', 21], ['block_cycles_cache_1', 17], ['block_cycles_cache_2', 41], ['block_cycles_cache_3', 86]]
[[['block_cycles_cache_0', 0], ['block_cycles_cache_1', 0], ['block_cycles_cache_2', 0], ['block_cycles_cache_3', 0]]
[[['block_cycles_cache_0', 1], ['block_cycles_cache_1', 1], ['block_cycles_cache_2', 1], ['block_cycles_cache_3', 1]]
[[['block_cycles_cache_0', 3], ['block_cycles_cache_1', 3], ['block_cycles_cache_2', 3], ['block_cycles_cache_3', 3]]
[[['block_cycles_cache_0', 1], ['block_cycles_cache_1', 1], ['block_cycles_cache_2', 1], ['block_cycles_cache_3', 1]]


opt constant in this round
fitness is -40875

iteration 1 begins

[[['block_cycles_cache_0', 1], ['block_cycles_cache_1', 1], ['block_cycles_cache_2', 1], ['block_cycles_cache_3', 1]]
[[['block_cycles_cache_0', 2], ['block_cycles_cache_1', 2], ['block_cycles_cache_2', 2], ['block_cycles_cache_3', 2]]
[[['block_cycles_cache_0', 21], ['block_cycles_cache_1', 17], ['block_cycles_cache_2', 46], ['block_cycles_cache_3', 86]]
[[['block_cycles_cache_0', 21], ['block_cycles_cache_1', 40], ['block_cycles_cache_2', 19], ['block_cycles_cache_3', 5]]
[[['block_cycles_cache_0', 9], ['block_cycles_cache_1', 59], ['block_cycles_cache_2', 22], ['block_cycles_cache_3', 11]]
[[['block_cycles_cache_0', 14], ['block_cycles_cache_1', 50], ['block_cycles_cache_2', 29], ['block_cycles_cache_3', 4]]
[[['block_cycles_cache_0', 7], ['block_cycles_cache_1', 59], ['block_cycles_cache_2', 39], ['block_cycles_cache_3', 19]]
[[['block_cycles_cache_0', 13], ['block_cycles_cache_1', 20], ['block_cycles_cache_2', 6], ['block_cycles_cache_3', 40]]


opt constant in this round is [['block_cycles_cache_0', 11], ['block_cycles_cache_1', 18], ['block_cycles_cache_2', 127], ['block_cycles_cache_3', 17]]
fitness is -40875
opt constant for gene algorithm is [['block_cycles_cache_0', 11], ['block_cycles_cache_1', 50], ['block_cycles_cache_2', 127], ['block_cycles_cache_3', 17]] fitness -40875
+ XiangShan git:(master) *
```

Design Space Exploration



```
block_cycles_cache_0 = 11
block_cycles_cache_1 = 12
block_cycles_cache_2 = 13
block_cycles_cache_3 = 15
```

Conf file





# Constantin

- **Usage:** Create signal
- XiangShan/coupledL2/src/main/scala/coupledL2/prefetch/TemporalPrefetch.scala

```
139 require(cacheParams.hartIds.size == 1)
140 val hartid = cacheParams.hartIds.head
141 // 0 / 1: whether to enable temporal prefetcher
142 private val enableTP = WireInit(Constantin.createRecord("enableTP"+hartid.toString, initialValue = 1.U))
143 // 0 ~ N: throttle cycles for each prefetch request
144 private val tpThrottleCycles = WireInit(Constantin.createRecord("tp_throttleCycles"+hartid.toString, initialValue = 4.U(3.W)))
145 // 0 / 1: whether request to set as trigger on meta hit
146 private val hitAsTrigger = WireInit(Constantin.createRecord("tp_hitAsTrigger"+hartid.toString, initialValue = 1.U))
147 // 1 ~ triggerQueueDepth: enqueue threshold for triggerQueue
148 private val triggerThres = WireInit(Constantin.createRecord("tp_triggerThres"+hartid.toString, initialValue = 1.U(3.W)))
149 // 1 ~ tpEntryMaxLen: record threshold for recorder and sender (storage size will not be affected)
150 private val recordThres = WireInit(Constantin.createRecord("tp_recordThres"+hartid.toString, initialValue = tpEntryMaxLen.U))
151 // 0 / 1: whether to train on vaddr
152 private val trainOnVaddr = WireInit(Constantin.createRecord("tp_trainOnVaddr"+hartid.toString, initialValue = 0.U))
153 // 0 / 1: whether to eliminate L1 prefetch request training
154 private val trainOnL1PF = WireInit(Constantin.createRecord("tp_trainOnL1PF"+hartid.toString, initialValue = 0.U))
```



# Constantin

- **Hands-on:** Pass constant via standard input stream

```
$ cd constantin           # now in tutorial/constantin
```

```
# Build emu with constantin (time consuming, use pre-built emu instead)
```

```
# bash step0-build.sh
```

```
# Run emu and pass constant
```

```
$ bash step1-basic.sh
```

```
please input total constant number
```

```
2
```

```
please input each constant ([constant name] [value])
```

```
tp_recordThres0 14
```

```
tp_throttleCycles0 3
```



- Hands-on: Pass constant via standard input stream

```
~/tu/test/t/constantin tutorial-new = ?4 > bash ./step1-basic.sh
emu-constantin compiled at Oct 28 2023, 17:38:43
please input total constant number
2
please input each constant ([constant name] [value])
tp_recordThres0 14
tp_throttleCycles0 3
Using simulated 32768B flash
isWriteICacheTable0 = 0
isWriteFetchToIBufferTable0 = 0
isWriteIfuWbToFtqTable0 = 0
isWritePrefetchPtrTable0 = 0
isWriteFTQTable0 = 0
isWriteBankConflictTable0 = 0
isWriteL1MissQMissTable0 = 0
isWriteLoadMissTable0 = 0
isFirstHitWrite0 = 0
isWriteLoadAccessTable0 = 0
isWriteL2TLbPrefetchTable0 = 0
isWriteL1TLbTable0 = 0
isWritePageCacheTable0 = 0
isWritePTWTable0 = 0
isWriteL2TLbMissQueueTable0 = 0
CorrectMissTrain0 = 0
isWriteInstInfoTable0 = 0
l1_stride_ratio0 = 2
l2_stride_ratio0 = 5
tp_trainOnL1PF0 = 0
tp_hitAsTrigger0 = 1
tp_triggerThres0 = 1
enableTP0 = 1
enableL3StreamPrefetch0 = 0
enableL1StreamPrefetcher0 = 1
nMaxPrefetchEntry0 = 14
always_update0 = 1
ForceWriteUpper_0 = 60
ForceWriteLower_0 = 55
tp_recordThres0 = 14
tp_throttleCycles0 = 3
tp_trainOnVaddr0 = 0
StoreWaitThreshold_0 = 0
ColdDownThreshold_0 = 12
enableDynamicPrefetcher0 = 1
StoreBufferThreshold_0 = 7
StoreBufferBase_0 = 4
l2DepthRatio0 = 2
l3DepthRatio0 = 3
depth0 = 32
Using simulated 8192MB RAM
The image is ./ready-to-run/linux.bin
The reference model is ./ready-to-run/riscv64-nemu-interpreter-so
The first instruction of core 0 has committed. Diff test enabled.
Core 0: EXCEEDING CYCLE/INSTR LIMIT at pc = 0x80000e4c
instrCnt = 1,000, cycleCnt = 10,235, IPC = 0.097704
Seed=0 Guest cycle spent: 10,239 (this will be different from cycle)
Host time spent: 44,688ms
```

```
please input total constant number
2
please input each constant ([constant name] [value])
tp_recordThres0 14
tp_throttleCycles0 3
```

```
ForceWriteUpper_0 = 60
ForceWriteLower_0 = 55
tp_recordThres0 = 14
tp_throttleCycles0 = 3
tp_trainOnVaddr0 = 0
StoreWaitThreshold_0 = 0
ColdDownThreshold_0 = 12
enableDynamicPrefetcher0 = 1
```



# Constantin

- **Feature:** Auto Solving of better constant values
- **Process**
  - Enable AutoSolving & prepare signals
  - Compile
  - Run basic demonstration
  - Set parameters for AutoSolving
  - Auto solve
  - Cleanup



# Constantin

- **Feature:** Auto Solving of better constant values

```
# Automatical parameter solver configuration file
```

```
$ cat constantin.json
```

- **Parameters defined**
  - properties of signals
  - target of Auto Solving
  - parameters of genetic algorithm
  - parameters of XS simulator



- Usage: Auto Solving of better constant values

*# Run auto solver, output will be the optimal constant currently found*

```
$ bash step2-solve.sh
```

*# The solver generates optimal parameters*

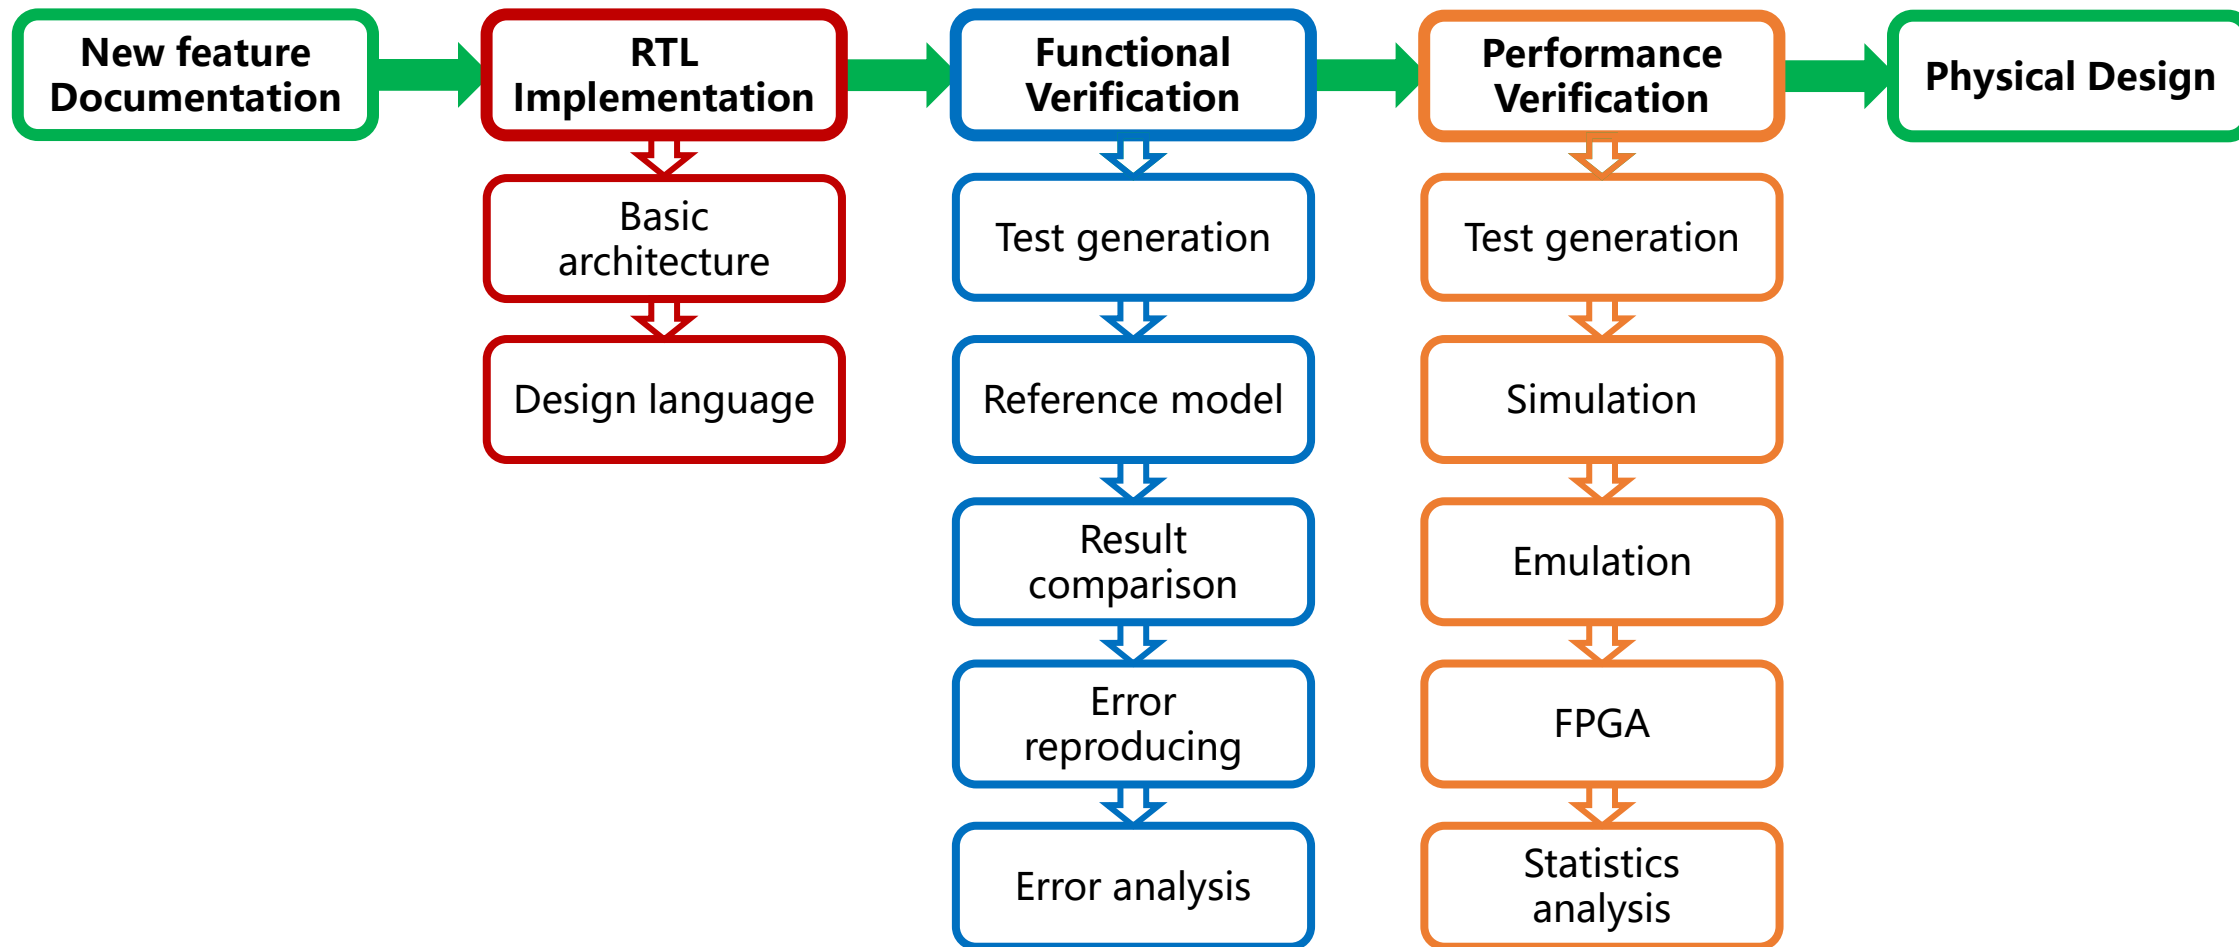
*opt constant for gene algrithom is*

```
['block_cycles_cache_0', XXX], ['block_cycles_cache_1', XXX],  
['block_cycles_cache_2', XXX], ['block_cycles_cache_3', XXX]] fitness 0
```





# Summary: MinJie Development Flows and Tools





北京开源芯片研究院  
BEIJING INSTITUTE OF OPEN SOURCE CHIP



中国开放指令生态（RISC-V）联盟  
China RISC-V Alliance

# Thanks!